

# CS2850 Operating System Lab

## Week 1: Introduction

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

References

System Programming

Key features of C

Compilers

Standard library

Control flow: `if`, `for`, `while` ...

# References

Brian W. Kernighan, Dennis Ritchie: [The C Programming Language](#) Prentice-Hall 1978 ISBN 0-13-110163-3

Randal Bryant, David O'Hallaron: [Computer Systems: A Programmer's Perspective C](#) Pearson Education Limited, 3rd edition, 2016 ISBN-13: 9781292101767.

[The GNU C Library Reference Manual](#)

# The Operating System (OS)

The OS is a layer of software that

- provides a better, simpler, cleaner, **model** of the computer and
- helps the user **handle resources**: processors, disks, printers, keyboard, display, ...

Two popular OS are UNIX and Windows.

# Why C-programming?

C is a **general-purpose** programming language. You can write almost anything in C.

C is **not tied** to any OS. Your programs will work on any machine.

**UNIX** is largely written in C.

# A relatively low-level language

C **does not** include

- × operators acting on **composite objects**, e.g. strings of characters, arrays or lists,
- × **Dynamical** memory allocation facilities,
- × **READ** or **WRITE** statements (you need to *call* dedicated functions),

# C is 'easy'

*“... keeping the language down to modest size has real benefits. Since C is relatively small, it can be described in a small place and learned quickly. A programmer can reasonably expect to know and understand and indeed regularly use the entire language”<sup>1</sup>*

---

<sup>1</sup>from Brian W. Kernighan, Dennis Ritchie: [The C Programming Language](#)

# ANSI C

C is machine-independent.

The program below works on computers with different OS.

```
#include <stdio.h>
int main() {
    printf("hello, world\n");
}
```

To run it, you need the *system-dependent* **executable**, a.out,

```
00000000: 01111111 01000101 01001100 01000110 00000010 00000001 .ELF..
00000006: 00000001 00000000 00000000 00000000 00000000 00000000 .....
0000000c: 00000000 00000000 00000000 00000000 00000011 00000000 .....
00000012: 00111110 00000000 00000001 00000000 00000000 00000000 >.....
....
```



# Compilation under UNIX

The OS produces a.out from the provided C code<sup>2</sup>.

To compile hello.c, use the **shell command**

```
gcc -Wall -Werror -Wpedantic hello.c
```

A useful **sanity check** of your program is run by entering<sup>3</sup>

```
valgrind ./a.out
```

---

<sup>2</sup>Use the additional option `-o yourExecutable` to change the executable name.

<sup>3</sup>valgrind is a powerful debugging tool for Linux programs.

# The Standard Library

ANSI C is based on a *established library* of **functions**.

You need standard library functions to

- **read or write** files,
- **allocate memory**,
- handle **strings**,
- ....

They are mostly written in C and may contain *a few* **non-portable** OS details, e.g. system call syntax.

# Headers

Write `# include <...>` to make a piece of the library accessible to your program.

`<stdio.h>`: input and output.

`<stdlib.h>`: memory allocation, process control.

`<unistd.h>`: system calls.

`<string.h>`: string-handling.

`<errno.h>`: error reporting.

`<math.h>`: common mathematical functions

# Control flow

The control flow fixes the order in which instructions are executed

The most used **control-flow statements** are

- **sequential instructions**: “;” (default line-by-line execution)
- **grouping symbols**: { ... }
- **selection commands**: if-else, switch, ...
- **repetition tools**: for, while, ...

## Example

A C program that prints “hello, world” several times

```
#include <stdio.h>
#define N 5
int main() {
    int i;
    for (i = 0; i < N; i++) {
        printf("%d) hello, world\n", i + 1);
    }
}
```

The output is

```
1) hello, world
2) hello, world
3) hello, world
4) hello, world
5) hello, world
```

# Notes

All variables need to be **declared** before using them.

`printf`<sup>4</sup> can print

- **simple strings**: `printf("hello, world\n");`
- **variable values**: `printf("%d \n", i);` where `%d` specifies that `i` should be printed as an **int**
- a **mix** of string and values:  
`printf("%d) hello, world \n", i);`

Add `/* .... */` and `// ...` to **comment** out multiple or single lines.

---

<sup>4</sup>Defined in `stdio.h`

# CS2850 Operating System Lab

## Week 2: Types, Variables, Functions

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

Variables

Functions

Types

Composite objects



# Programs

C programs consist of **functions** and **variables**

A function is a set of **statements** that specify the operations to be performed

A variable is a **location in storage** with three attributes:

- an **identifier** or *name*,
- a **storage class** determining the variable *lifetime*, and
- a **type** determining the *meaning* of the value found in the identified storage

# Declaration

Functions and variables should always be **declared** before they are used, i.e. you need to write

```
storage-class variable-type variable-name;
```

The **lifetime** of a variable is specified by *where* it is declared (and by its *storage class*).

The declaration of a variable **does not initialize** it automatically.<sup>1</sup>

---

<sup>1</sup>But you can separate their declaration and definition.

## Example

```
#include <stdio.h>                                1
void printInt(int j);                              2
int main() {                                       3
    int i;                                        4
    i = 1;                                        5
    printInt(i);                                  6
}                                                  7
void printInt(int k) {                             8
    printf("i=%d\n", k);                          9
}                                                  10
```

# Types, names, storage classes

The variable *type* specifies how the stored bytes should be *interpreted*.

Variable **names** are sequences of letters and digits used for referring an object, i.e. `i` or `printInt`.

There are two important storage classes<sup>2</sup>:

- **automatic**: the variable is local to the block and reinitialized every time the function is called and
- **static**: the variable keeps its value across different function calls.

---

<sup>2</sup>By default, *local* variables are automatic, *global* variables are static.

# Functions

A function is **defined** by writing

```
return_type  function_name(arguments) {statements};
```

where

- `return_type` is the *type* of the return value,
- `function_name` is the function name (used to *call* it,
- `arguments` denotes the *name and type* of the function parameters, and
- `statements` is the set of *operations* to be performed.

# Function calls

A function is **called** by writing its name

```
f(argument_value_1, argument_value_2, ... );
```

where `argument_value_I` is the value of parameter `I` (at the right place).

In the function declaration, parameter names are *arbitrary*. Only their *type* is **strictly required** (see the code above for an example).

# Types

There are a few **basic** data types:

- **char** a single character, e.g. a
- **int** an integer, e.g. -1234
- **unsigned int** a positive integer, e.g. 12345
- **float** a single-precision floating point, e.g. 23, 923, and
- **void**, an empty set of values <sup>3</sup>.

The operator **sizeof()** returns the (architecture dependent) number of bytes of any given basic type.

---

<sup>3</sup>void is not a regular type, e.g. you cannot declare a void variable.

## Example

Run this program to check the size of the basic C types on your system

```
#include <stdio.h>
int main() {
    printf("sizeof(char)=%lu\n", sizeof(char));
    printf("sizeof(int)=%lu\n", sizeof(int));
    printf("sizeof(unsigned int)=%lu\n", sizeof(unsigned
        int));
    printf("sizeof(float)=%lu\n", sizeof(float));
}
```



## More on int and char

**Integer constants** can be used to initialize a variable and given in binary (e.g. `int a = 0b01010101;`), decimal (`int a = 85;`), hexadecimal (`int a = 0x55;`) or octal (`int a = 0125;`).

**Character constants** are characters enclosed in single quotes, (e.g. `char a = 'x';` or `char a = '\n';`).

The numerical value of a **char** is the **ASCII** integer code corresponding to it, e.g. check the output of

```
printf("ASCII(r)=%d\n", 'r');
```

Printable characters, e.g. `r`, `$`, or `3`, are always positive but plain char's can be **signed or unsigned** (depending on the machine).

## Example

```
#include <stdio.h>
int main() {
    char a = -123;
    printf("character=%c\n", (char) a);
    printf("signed int=%d\n", (int) a);
    printf("signed int (octal)=%o\n", (int) a);
    printf("unsigned int=%u\n", (unsigned int) a);
    printf("unsigned int (octal)=%o\n", (unsigned int) a);
}
```

Type casting specifiers as `(unsigned int)` force a *specified interpretation* of the following variable.

While `-123` does not correspond to a *valid character*, the *stored bytes* (printed in octal) can be interpreted in different ways.

# Operators

**Arithmetic** operators: +, -, \*, /, and %.

**Relational** operators: ==, !=, >, >=, <, and <=.

**Logical** operators: && (and), and || (or).

The **negation operator**, “!”, converts a non-zero operand into a 0, and a zero operand into 1 i.e. `if (!operand)` and `if(operand == 0)` are *equivalent*.

% is the **modulus operator**, e.g.  $1\%2 = 1$ ,  $2\%1 = 0$ ,  $3\%2 = 1$ , and  $2\%3 = 2$ .

# Arithmetics

Operators may cause **conversion** of the value of an operand from one type to another.

For example, **integer division** may or may not **truncate** any fractional part, e.g.  $3/2 = 1$  but  $3/2. = 1.5$  and  $3./2 = 1.5$ .

The full list of **arithmetic conversion rules** is in Appendix A6.5 of The C Programming Language.

# Non-basic types

There is a *conceptually infinite* class of **derived types**, which are built from the basic types in various ways.

The most important *derived types* are

- **strings**: null-terminated lists of char's,
- **arrays**: lists of objects of a given type,
- **functions**: sets of statements returning objects of a given type,
- **pointers**: memory addresses of objects of a given type, and
- **structures**: general composite objects containing objects of different types.

# The size of composite objects

The **structure** of composite objects, e.g. the number of entries of an array, should be declared before they are used and *cannot be changed at run time*.

The **size** in bytes of composite objects can be obtained using **sizeof**

```
printf("sizeof(char[10])=%lu\n", sizeof(char[10]));  
printf("10 * sizeof(char)=%lu\n", 10 * sizeof(char));
```

# Strings (1)

The program in the next slide,

- i) *declares and initialises* a string (constant),
- ii) *prints* the string using `printf` and the format specifier `%s`,
- iii) prints a specific *character* of the string and non-initialised value *outside* the string.

## Strings (2)

```
#include <stdio.h> 1
int main() { 2
    char *s = "hello, world\n"; 3
    printf("s=%s", s); 4
    printf("s[7]=%c\n", s[7]); 5
    printf("s[100]=%c\n", s[100]); 6
} 7
```



## Strings (3)

The program *does not crash* if you try to access **uninitialised entries** of `s` but the value stored there is *unpredictable*

The **format specifier** `%s` allows you to use `printf` for printing `s` with `as` as a unit.

The program *knows where the string terminates* because strings are **null-terminated** lists of `char`'s, i.e. the last `char` of a string is *always* a `\0`.

## Arrays (1)

The program in the next slide,

- i) *declares* an array of 10 `int`, i.e. allocate the memory space to store 10 `int`,
- ii) *loads* random integers to its entries, through component-wise assignments, and
- iii) prints the vector components with `printf` and the format specifier `%d`.

## Arrays (2)

```
#include <stdio.h>
#include <stdlib.h>
void loadVector(int a[], int size) {
    for (int i=0; i<size; i++) a[i] = ((float) 10 * rand
        ())/RAND_MAX;
}
void printVector(int a[], int size) {
    for (int i=0; i<size; i++) printf("a[%d]=%d\n", i, a[i
    ]);
}
int main() {
    int size = 10;
    int a[size];
    loadVector(a, size);
    printVector(a, size);
    printf("a[100]=%d\n", a[100]);
}
```

## Arrays (3)

If you try to access uninitialized entries of `a` but the value stored there is *unpredictable*.

The program *let you access* uninitialized memory regions, i.e. `a[100]`.

You need a customized function to load and print `a` as a single unit.

`loadVector` and `printVector` do not know the length of the input because `a` is passed as *the address of `a[0]`*.

# CS2850 Operating System Lab

## Week 3: Memory basics, Pointers, Arrays

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

Memory, pointers, arrays, and strings

Address arithmetic

Arrays and functions

Pointers to pointers and command-line arguments

# Memory basics

The memory is a *large array* of **memory cells** (bytes).

The array contains the **entire state** of your programs: variables, constants, data, and machine code.

Data are stored at specific **memory addresses**.

**Pointers** are *variables* for storing the memory address of other variables.



## Two useful operators

```
int i = 1;  
int *ip = &i;
```

The *address operator*, **&**, returns the *address of* *i*.

The *dereferencing operator*, **\***, returns the value stored at a given address.

*ip* is a pointer variable storing the address of *i*.



## Pointers at work

```
int i = 1;
int *ip, *iq;
ip = &i;
*ip = *ip + 1;
iq = ip;
```

You can use `*ip` instead of `i` in **any context**, e.g. `*ip = *ip + 1`; adds 1 to `i`.

You can *initialize* or *redirect* pointers by copying the content of **another pointer of the same type** to it, e.g. `iq = ip`;

# Types and pointers

Pointers store the address of specific **data type**, e.g. `int *ip;` says that `ip` is a *pointer to integers*, i.e. `*ip` is a `int`.

All pointers have the **same size**, 8 bytes.

The *generic pointer type*, `void *`, can be cast to `void *` and back.

## Example

This program returns the content of a `void*` address that is *casted* to a pointer to `int`.

```
int main() {  
    int i = 1;  
    int *ip = &i;  
    void *iq;  
    iq = (void *)ip;  
    return *(int *)iq;  
}
```

To see the return value on the terminal, run

```
gcc -Wall -Werror -Wpedantic pointers.c  
./a.out  
echo $?
```

# Declaring arrays

The following declaration **allocates** 10 *consecutive* blocks of 4 bytes named `a[0]`, `a[1]`, ..., `a[9]`,

```
int a[10];
```

As `a[0]`, `a[1]`, ... are *all* integers the program only needs to know the address of the **first element**, i.e. the *pointer and type to the first element*.

# Pointers and arrays

The following defines a **pointer** to the first element of a,

```
int a[10];  
int *pa = &a[0];
```



Pointers and arrays are *closely related*: the **value** of a (without brackets) is the *address of its first element*.

# Strings

```
char *s = "hello world";
```

**Strings** are *null-terminated* arrays of `char`, i.e. their last char is `'\0'`.

The **null-termination** lets the program find the *end of the string*.

There are *no C operators for processing strings as units*. But you can use `printf("s=%s\n", s);` to print `s` or `as` because they are **null-terminated**.

## String constants and character arrays are different

```
char *s = "string constant";  
char as[20] = "character array";
```


s is a **pointer** *to a constant* (un-modifiable) piece of memory.

as is the **address** of a 20-byte *character array* (you can write on it).

# Address arithmetics

Let **a** be an array of 10 int and **pa** a pointer to int.

After writing **pa = a;**, **pa + 1** points to **a[1]** and **pa + 4** points to **a[4]**<sup>1</sup>



**a[i]** and **\*(a+i)** refer to the same object (the content of **a[i]**).

---

<sup>1</sup>The **value** of **a** (without brackets) is the *address of its first element*



## Note

**Portions** of `s` (or `as`) can be accessed by specifying the address of a single character within them.

The following lines print the **substrings** "constant" and "array",

```
char *s = "string constant";  
char as[20] = "character array";  
printf("%s", &s[7]);  
printf("%s", s + 7);  
printf("%s", &as[10]);  
printf("%s", as + 10);
```

# Pointers and functions

Arguments are passed to functions **by value**.

Functions **cannot modify** a variable *defined in the calling function*.

To *save the changes* you need to

- define a function with **pointer arguments**, e.g.

```
void f(int *a) { *a = 5; }
```

- **pass a pointer** when you call f.

```
int a = 3;  
f(&a);  
printf("a=%d\n", a);
```

# Pointers to pointers

Pointers can store the address of other pointers.

You can use an **array of pointers** to char to store a *list of strings*,

```
char *sa[10];
sa[0] = "hello";
sa[1] = ", ";
sa[2] = "world";
sa[3] = "!";
sa[4] = NULL;
int i = 0;
while (*(sa + i)) {
    printf("sa[%d]=%s\n", i, sa[i]);
    i++;
}
```

# Command-line arguments

C programs accept **command-line** arguments through a **strings array** called `argv`.

```
int main(int argc, char **argv) {  
    int i = 1;  
    while (i < argc) {  
        printf("argv[%d]=%s\n", i, argv[i]);  
        i++;  
    }  
}
```

The output is as before,

```
cim-ts-node-02$ ./a.out hello , world !  
argv[1]=hello  
argv[2]=,  
argv[3]=world  
argv[4]=!
```

# CS2850 Operating System Lab

## Week 4: processes and IO

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# outline

System calls

File system

Processes

fork

# System calls

The **operating system** provides its services through a set *system calls*.

Accessing to the OS is mainly needed for

- **input/output**,
- **file** handling,
- **memory** allocation,
- ...

# A portable interface

Different OS's may offer similar services in *different ways*.

In a C program, system calls are **architecture-independent**, e.g. `printf` performs a system call and its syntax is machine independent.

The **portable interface** is a set of *standard library functions*.

These routines *depend on the host OS*, i.e. they are written in terms of the facilities of a *specific OS*.



# System calls from a C program

From a C program, you can make

- Task-dedicated **UNIX system calls**, e.g.

```
int n_read = read(int fd, char *buf, int n);
```

- **Direct system calls**, through a general function declared in `unistd.h`. (for system calls that have no *wrapper* in the standard library), e.g.

```
syscall(3, ... );
```

where 3 is the number of the *read* system call.

# syscall

syscall is a flexible function with a variable number of parameters,

```
long int sys_return = syscall(long int SYS_NO, ... );
```

- SYS\_NO is the **system call number**, which identifies each kind of system call,
- the remaining *parameters* depend on the system call and the architecture,
- the **return value** is the return value of the specific system call.

syscall is useful for system calls that have **no wrapper** in the C library.

## Example

The following program prints hello, world on screen.

```
#include <unistd.h>
int main() {
    syscall(1, 1, "hello world \n", 20);
}
```

What is the meaning of the first two arguments? What happens if you change the last argument to 6?

## Low-level input-output

You can use `syscall` to implement low-level IO operations.

```
n_read = syscall(SYS_read, fd, buf, n)
n_written = syscall(SYS_write, fd, buf, n)
```

where `SYS_read` and `SYS_write` are the numbers associated with the *read* and *write* system calls in `sys/syscall.h`.

# Primitive input-output

*Reading and writing* is easier with

```
int n_read = read(int fd, char *buf, int n);  
int n_written = write(int fd, char *buf, int n);
```

- `n_read` or `n_written` is the number of bytes actually transferred (if an error occurs, `n_read < n`, idem `n_write`),
- `fd` is a *file descriptor* (see below),
- `char *buf` is a character array,
- `n` is the number of bytes to be transferred.

`read` and `write` appear in *higher-level* functions that read or write, e.g. `getchar` or `printf`.

## Example (1)

This program copies its input into its output.

```
#include <unistd.h>
int main() {
    char buf[20];
    int n;
    while((n = read(STDIN_FILENO, buf, 20)) > 0)
        write(STDOUT_FILENO, buf, n);
}
```

The **file descriptors** `STDIN_FILENO` and `STDOUT_FILENO` are defined in `unistd.h` and refer to *standard input* (the keyboard) and *standard output* (the terminal).

## Example (2)

This program is equivalent to the previous one.

```
#include <unistd.h>
int main() {
    char buf[20];
    int n;
    while((n = syscall(0, 0, buf, 20)) > 0)
        syscall(1, 1, buf, n);
}
```

Read more about read and write on [this page](#) and more about syscall on [this page](#) of the [The GNU C Library Reference Manual](#).

## syscall codes

Run the following program to print on screen the syscall codes for a few IO and process control operations.

```
#include <stdio.h>
#include <sys/syscall.h>
int main() {
    printf("SYS_read=%d\n", SYS_read);
    printf("SYS_write=%d\n", SYS_write);
    printf("SYS_open=%d\n", SYS_open);
    printf("SYS_close=%d\n", SYS_close);
    printf("SYS_fork=%d\n", SYS_fork);
    printf("SYS_getpid=%d\n", SYS_getpid);
}
```

Why do you need to include `stdio.h` here? Why can you use file-IO facilities to read from and write on the terminal?



# Streams

A **file** is a *stream of data*, i.e. a sequence of bytes.

What happens when you **open** a file?

1. the **kernel** checks your right to do so (Does the file exist? Do you have permission to access it?),
2. the kernel returns to the program a small non-negative integer called a **file descriptor**, and
3. the kernel keeps track of all information about the open file, e.g. a file **position** (offset from the beginning of a file).

# File references

*Direct system calls*, e.g. `read`, `write`, and `syscall`, refers files through their **file descriptor**.

*Higher-level IO facilities* defined in the standard library refer to files through their **file handle**.

A file handle is an object of type `struct` containing

- a pointer to a **buffer**,
- the number of **characters left** in the buffer,
- a pointer to the **next position** in the buffer,
- the **file descriptor**,
- various **flags**, e.g. describing the read/write *mode*,
- ...

# Unix I/O

In a C program, the easiest way to open a file is to call

```
FILE *fopen(char *name, char *mode);
```

- FILE is the **file handle-type**.
- The name parameter is the **file name** (a string constant).

The mode parameter controls how the file is to be opened ("r" for *read-access*, "w" for *write-access* and "a" for *append access*).

# File descriptors and file handles

*File handles*, i.e. struct of type FILE, are composite objects and contain all the information you need to access an open file.

In most cases, file handles are *passed to functions* by passing a **pointer to FILE**.

*File descriptors* are simple integers. You can obtain the file descriptor of an open file from the file handle of an open file with

```
int fd = fileno(FILE *fh);
```

Normally, IO functions defined in `stdio.h`, e.g. `fopen` and `fileno`, refer to open files through file handles.

## Lower-level file handling

The *lower-level* versions of `fopen` are defined in `fcntl.h` and return a **file descriptor**.

```
int open(char *name, int flag, int perms)
int creat(char *name, int perms)
```

See [The GNU C Library Reference Manual](#) for more details about the arguments `flags`, a parameter that controls how the file is opened, and `perms`, a number encoding the user's permissions.

## stdin and stdout

In UNIX, all **IO** operations are done by reading or writing files all peripheral devices, even the keyboard and the screen, are **files** in the file system.

Printing on the **terminal** or reading from the **keyboard** is a special case.

All processes automatically open three files called `stdin` (*standard input*), `stdout` (*standard output*), and `stderr` (*standard error*), with **fixed file descriptors 0, 1, and 2**.

`stdin`, `stdout`, and `stderr` are **file handles** defined in `stdio.h`.

## High-level reading and writing

In a C program, the easiest way to read and write on file is to use

```
int fscanf(FILE *f, char *format, ...)
int fprintf(FILE *f, char *format, ...)
```

Example: the following two calls are **equivalent**:

```
fprintf(stdout, "hello world");
printf("hello world");
```

Read more about `fscanf` and `fprintf` on the [The GNU C Library Reference Manual](#).

# Processes

A process is the OS **abstraction** of a program in execution.

Each program runs in the **context** of some process.

The context includes all **state information** needed to run the process: program code, data, stack, program counter, environment variables, . . . .



# Concurrent processes

You always have the *illusion* that your program is the **only one** running on the system.

This happens because a process has

- an *independent* **control flow** and
- a *private* **address space**, i.e. the memory associated with a process cannot be read or written by any other process.

*Multitasking* is needed for running multiple *concurrent* processes in an apparently simultaneous way.

# Process handling

Unix provides the basic **system calls** for manipulating processes from C programs.

A process has a unique **process identifier** (PID) obtained by calling

```
pid_t getpid(void);
```

defined in `unistd.h`.

The type of the return value (a positive integer), `pid_t`, is defined in `types.h`.

You can print the process identifier of your running program at any time by writing

```
printf("pid=%d\n", getpid());
```

# States of a process

A process can be in three **states**:

- i) *running*: **executing** on the CPU,
- ii) *stopped*: the execution is **suspended** (until it receives a signal),  
or
- iii) *terminated*: the process is stopped **permanently**

## Child processes

A process *can create a new running process* by calling

```
pid = fork();
```

which is defined in `unistd.h` and **returns twice**:

1.  $pid = 0$  in the **child** process, and
2.  $pid = pid\_child$ , where `pid_child` is the *pid of the generated child process* in the **parent** process.

`fork` returns  $-1$  if the process creation failed.

# Process termination

A process can be **terminated** by writing (in main)

- **return** *i*;

which *returning the integer i*,

- `exit()`;

which has no return value.

Note: you need to include `stdlib.h` to use `exit`.

## Child and parent are quasi-identical

`fork` creates a child process that is **almost identical** to its parent.

The child gets a *separate copy* of the parent's (user level) **address space**: code, data, heap, user stack.

The child gets copies of all parent's **open file** descriptors, e.g. the child can read and write any files that are *open in the parent when `fork` is called*.

This is why they can both print on `stdout`.

## Parent and child are distinct

The parent and the child have **different PID's**.

`fork` is called once by the parent, but it **returns twice**, once in the parent and once in the child.

The parent and the child processes **run concurrently** changes made after calling `fork` are **private** i.e not reflected in the other process.

Read more about `fork` on [The GNU C Library Reference Manual](#).

# Waiting for a child process to terminate

Unix provides a few more **system calls** for manipulating processes from C programs.

A parent process **waits** for a child to terminate by calling the

```
pid_t wait(int *status)
```

which is defined in `sys/wait.h`.

The parameter `status` can be used to encode some *information about the termination of the child process*,<sup>1</sup>, e.g. its return value.

The return value is *the PID of the child that terminates* or `-1` if the process has no children

---

<sup>1</sup>Set `&status` to `NULL` if you do not need this information



## Suspending a process

A process can be **suspended** by calling

```
unsigned int sleep(unsigned int secs)
```

which is defined in `unistd.h`.

`secs` is the number of seconds the process will sleep.


The return value is the number of seconds left to sleep (if it returns prematurely).

We suggest you read more about `wait` and `sleep` on [The GNU C Library Reference Manual](#).

# Examples (1)


```
#include <stdio.h>
#include <unistd.h>

int main(){
    fork();
    printf("hello\n");
}
```



```
#include <stdio.h>
#include <unistd.h>

int main(){
    fork();
    fork();
    printf("hello\n");
}
```



## Examples (2)

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main() {
    if (!fork())
        printf("the PID of the child is %d\n", getpid());
    else{
        printf("the return value of wait(NULL) is %d\n",
            wait(NULL));
        printf("the PID of the parent is %d\n", getpid());
    }
}
```

The program above prints

```
the PID of the child is 3384654
the return value of wait(NULL) is 3384654
the PID of the parent is 3384653
```

# CS2850 Operating System Lab

## Week 5: the UNIX shell

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

How the shell works.

The UNIX shell

Wild cards

Bash programming

# How the shell works

A *shell* is a **user-level process** for interacting with your machine.

The process *starts* when you open the terminal, e.g. when you log in using puTTY.

The shell prints a **prompt** and waits for you to *type a command* in the terminal and *press enter*.

The shell **executes the instruction** and prints a new prompt when it is ready to *accept a new command*.

# Command execution

When you run a command the shell creates a **new child process**.

The child process **executes** the corresponding program, e.g. `ls` is a program name. <sup>1</sup>

The shell **waits** for the child process to terminate.

After the shell has reaped the child process, it prints a prompt again and waits for the **next input**.

---

<sup>1</sup>The shell assumes the first word of the input is an executable file, like `./a.out`.

## Example

The user moves to Desktop/operatingSystems/, looks at the directory content, and prints their location.

```
cd Desktop/operatingSystems/  
ls  
CS2850Labs  intros  
pwd  
/home/cim/staff/ugqm002/Desktop/operatingSystems/  
ls -l CS2850Labs/  
total 72  
-rwx----- 1 ugqm002 staff 16832 Sep 21 17:56 a.out  
drwx----- 2 ugqm002 staff  8192 Oct  4 08:45 example  
drwx----- 2 ugqm002 staff  8192 Sep 23 13:01 week1  
drwx----- 3 ugqm002 staff  8192 Sep 30 17:23 week2  
drwx----- 2 ugqm002 staff  8192 Oct  7 18:08 week3
```



# fork and execve

Child processes are created with `fork` and used to invoke a `loader` (`execve`) that

1. `replaces` the child's virtual memory with the info needed *to execute the command*, e.g. the `code` of the command entered by the user,
2. `runs` the specified program (within the child process), and
3. `terminate` the child process.

## Example

This program works like a shell and prints `hello, world!` on the screen.

```
#include <unistd.h>
#include <wait.h>
int main() {
    char *argv[2];
    argv[0] = "a.out";
    argv[1] = NULL;
    if (fork() == 0) execv(argv[0], argv);
    wait(NULL);
}
```

`a.out` is the executable of the following code

```
#include <stdio.h>
int main() {
    printf("hello, world!\n");
}
```

# The UNIX shell

The **first** Unix shell was developed for Unix by Steven R. Bourne in 1974.

The *Bourne shell* is now the standard shell in UNIX systems.

Each process created by a Unix shell has three **open files**:

1. standard input (`fd = STDIN_FILENO = 0`),
2. standard output (`fd = STDOUT_FILENO = 1`), and
3. standard error (`fd = STDERR_FILENO = 2`).

## Extra features

IO **redirection** operators, `<` and `>`.

**Example:** the following lines *write the output* of `ls -l` on `output.txt` and print on the screen the first lines of `output.txt`.

```
ls > output.txt  
head < output.txt
```

A **pipe** command, `|`, to send a command output to another command.<sup>2</sup>

**Example:** the following line prints only the files and directories modified on October 25th.

```
more output.txt | grep "Oct 25"
```

---

<sup>2</sup>Without generating extra files.

# Shell commands

Program	Typical Use
<code>cat</code>	Concatenate files to standard output
<code>chmod</code>	Change file protection mode
<code>cp</code>	Copy one or more files
<code>cut</code>	Cut columns of text from a file
<code>grep</code>	Search a file for some pattern
<code>head</code>	Extract the first lines of a file
<code>ls</code>	List directory
<code>make</code>	Compile files to build a binary
<code>mkdir</code>	Make a directory
<code>ps</code>	List running processes
<code>rm</code>	Remove one or more files
<code>rmdir</code>	Remove a directory
<code>sort</code>	Sort lines alphabetically
<code>tail</code>	Extract the last lines of a file
<code>tr</code>	Translate between character sets

# ls-l

The **full information** about the files in a directory is obtained by running

```
ls -l directoryPath
```

and includes the *location*, *type*, *size*, the owner and owner's group, their *permissions*, and the last time the file was *modified*, e.g. <sup>3</sup>

```
ls -l
total 96
-rwx----- 1 ugqm002 staff 16840 Oct 20 11:08 a.out
-rw----- 1 ugqm002 staff    63 Oct 20 12:11 helloWorld.c
-rw----- 1 ugqm002 staff   496 Oct 20 12:27 output.txt
drwx----- 2 ugqm002 staff  8192 Sep 10 17:18 wl
```

---

<sup>3</sup>Directories have a 'd' as first character.

# Permissions

The permissions of files are encoded in 10-character strings, specifying the permissions of the *user*, the user *group*, and the *rest of the world* to *read*, *write*, and *execute* the file (in the order).

The 10 characters below show that the user can read, write, and execute, the user group can read and execute, and the rest of the world can only execute. e.g.

```
-rwxr-x--x
```

## chmod

You can **set or change** the permission bits with

```
chmod [who] operator [permissions] file_name
```

where  $\text{who} \in \{\text{u}, \text{g}, \text{o}, \text{a}\}$ ,  $\text{operator} \in \{+, -, =\}$ , and  $\text{permission} \in \{\text{r}, \text{w}, \text{x}\}$ .

**Example:** If you want to make a file **executable to yourself** you need

```
chmod u+x file_name
```

Check Section 1.4 of [Linux and Shell Programming](#) for using `chmod` in **absolute mode**, where the operation string with an octal number, e.g. 0100 means *"the owner can execute"*.



# Shell variables

The shell uses **environment variables** to set a working environment when you log in.

You can inspect *all* environment variables by running **env**, e.g.

```
env
SHELL=/bin/bash
LANGUAGE=en_GB:en
PWD=/home/cim/staff/ugqm002/Desktop/operatingSystems
    /2021/labs/week5
...
```

You can print a list of **selected variables**, VARNAME1, VARNAME2, ..., by running

```
echo $VARNAME1 $VARNAME2 ...
```

# Setting variable values

You can set the value of VARNAME to a specific value using

```
VARNAME=value
```

This changes the variable only for the **current process**<sup>4</sup>

You can clear a variable with

```
unset VARNAME
```

---

<sup>4</sup>Use `export VARNAME` to affects its **subprocesses** (but not its parent process).

## Example

Try the following commands to understand how this works

```
ONE=1; export ONE; TWO=2
```

```
./changeVariables.sh
```

```
1
```

```
1
```

```
3
```

```
echo $ONE
```

```
1
```

```
echo $TWO
```

```
2
```

`changeVariables.sh` is a file containing the following lines

```
#!/bin/bash
```

```
echo $ONE; echo $TWO
```

```
ONE=1; TWO=3; export TWO
```

```
echo $ONE; echo $TWO
```

# Wild cards

The shell has a set of **pattern-matching meta-characters** to match strings based on patterns:

- `*` matches any string including a null string, e.g. `ls *.c`,

- `?` matches any one character, e.g. `ls ??Script.sh`,

- `[...]` matches any characters enclosed in the square brackets, e.g. `ls *[S]*`, and

- `[!...]` matches any characters other than the characters following the exclamation mark, e.g. `ls *. [!s]*`

The resulting **general expressions** can be used for running shell commands.<sup>5</sup>

---

<sup>5</sup>The shell expands all wild cards before executing the commands.

## Example

The following bash lines combine shell commands and regular expressions

```
ls -l *.* > output.txt  
grep '\.c' output.txt
```

`ls -l *.* > output.txt` writes the info of all items with name matching `item_name.ext_name` on `output.txt`

`grep '\.c' output.txt` picks the lines corresponding to all `.c` files.

See more about `grep` (global regular expression print) in Chapter 8 of [Linux and Shell Programming](#).

# Shell scripts

You can compose a **list of commands** for the shell to execute.

A **shell script** is a file holding *a program built with shell commands*.

Scripts can use some of the *usual constructs* to make simple **decisions** and **loops**, i.e. the shell is a **programming language**.

## Example

```
#!/bin/bash
i=start; echo $i > output.txt
for i in 1 2 3 4 5 6 7 8
do
    echo $i "hello world" >> output.txt
done
more output.txt
```

The script above prints the following lines.

```
./myScript.sh
start
1 hello world
2 hello world
...
7 hello world
8 hello world
```

# Script variables

**Local variables** can be declared, modified, and printed on the terminal, e.g.

```
VARNAME=value  
$VARNAME=newValue  
echo "VARNAME=" $VARNAME
```

**Special variables** can be used to process the program input, e.g.

\$#: number of parameters passed to the script,

\$0: script name,

\$1: first parameter,

\$2: second parameter

...

\$?: exit value of last command.



# Tests

Let VAR be a *string* and N an *integer*, e.g. VAR=one and N=1.

You can make **tests** through

```
if [ $VAR = test_string ]  
if [ $N -eq test_value ]
```

**Note:** all *empty spaces* are strictly required.

You can modify **numerical variables** using `expr`, e.g.

```
N=`expr $N + $N + 1`
```

# Loops

You can use **for-loops** and **while-loops** as in many other languages, e.g.

```
#!/bin/bash
VAR=0
echo "VAR=" $VAR
for i in 1 2 3; do
    VAR=`expr $VAR + $i`
done
echo "VAR=" $VAR
while [ $VAR -ge 0 ]; do
    VAR=`expr $VAR - 1`
done
echo "VAR=" $VAR
```

# Example

Run this script with 2 inputs,  $\text{string} \in \{\text{plus}, \text{minus}\}$  and  $\text{value} \in \mathbb{N}$

```
#!/bin/bash
VAR=$1; N=$2
echo "VAR=" $VAR ", N=" $N
if [ $VAR = plus ]; then
    N=`expr $N + 1`; VAR=good
else
    if [ $N -ge 0 ]; then
        N=`expr -1`
    fi
    N=`expr $N - 1`; VAR=...
fi
echo "VAR=" $VAR ", N=" $N
if [ $N -lt 0 ]; then
    echo "wait..."
    while [ $N -le 0 ]; do
        N=`expr $N + 1`; echo "N=" $N
    done
    VAR=good
fi
echo "VAR=" $VAR ", N=" $N
```

# CS2850 Operating System Lab

## Week 6: Dynamic memory

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# outline

The heap and the stack

Static memory allocation

Dynamic memory allocation

# Processes


A process is a *running program*.

Each process has a private **address space**, i.e. a set of memory locations *labeled* from 0 to some maximum.

The address space contains the program

- *executable file*,
- *data*,
- *stack*, and
- *heap*.

# Virtual address space



The process *virtual* address space is a list of memory locations.

The program reads from and writes to both the **heap** and the **stack**.

# The stack

The stack is the region *on the top* of the virtual address space used to implement *function calls*.

It *expands and contracts* when the program runs, e.g.

- it **grows** each time a function is called, e.g. to allocate *local variables*, and
- it **contracts** when a function returns as all local variables are *discarded*.



# The heap

The heap is a *collection* of various-sized blocks, each block being a slot of virtual memory that is either **allocated** or **free**.

The code and data areas are *just below* the run-time heap.

The heap *expands and contracts* dynamically at run time as **dynamic memory** is allocated.

The **kernel** maintains a variable that points to the *top of the heap*.

## Example

Run this program to check that *static* and *dynamic* objects are allocated in different address space regions.

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    char a = 'a';
    char b[10];
    char *aDyn=malloc(sizeof(char));
    *aDyn = 'A';
    char *bDyn=malloc(10 * sizeof(char));
    for (int i=0; i<10; i++) {
        *(b + i) = 'b';
        *(bDyn + i) = 'B';
    }
    printf("a=%c and &a=%p\n", a, (void *) &a);
    printf("b[0]=%c and &b[0]=%p\n", *b, (void *) b);
    printf("**aDyn=%c and &p[0]=%p\n", *aDyn, (void *) aDyn);
    printf("bDyn[0]=%c and &bDyn[0]=%p\n", *bDyn, (void *) bDyn);
    free(aDyn);
    free(bDyn);
}
```

# Static memory allocation

In C, the amount of memory used to allocate variables in the stack is determined at **compile time**.

Stack variables are declared by writing

```
char a;  
int b[10];
```

The size of a stack variable can not be changed *at runtime*. In other words, the memory allocation is **static**.

# Dynamic memory allocation

The standard library provides functions for allocating memory in the run-time **heap**.

Heap memory blocks are allocated **dynamically** and it is possible to change the size of the corresponding variables.

To change the size of a heap-allocated variable, you need to **call a library function** that *finds* a new suitable slot and *frees* the old one.

# Memory allocator

The library function for allocating memory in the heap is

```
void *malloc(size_t size)
```

`size_t` is the size of the block to be allocated.

For example, you can allocate a *memory region* for hosting an array of 10 integers with

```
int *a;  
a = malloc(10 * sizeof(int))
```

# malloc

malloc is defined `stdlib.h` and

- returns a void pointer <sup>1</sup>, if the allocation is successful, and NULL otherwise, and
- creates a new **un-typed** and **uninitialised** block of memory of size `size`.

After allocating the right amount of memory you can cast the void pointer to **any required type**.

---

<sup>1</sup>A pointer to **unspecified** data

# How not to write your program

*“ISO C99 introduced the capability to have array dimensions be expressions that are computed as the array is being allocated, and recent versions of gcc support most of the conventions for variable-sized arrays in ISO C99.<sup>2</sup>*

**In this course, you should ignore these extensions and *always* use `malloc` to allocate memory dynamically.**

---

<sup>2</sup>from Section 3.8.5 of [Computer Systems: A Programmer's Perspective C](#)

## realloc and free

Other memory-allocation functions defined in `stdlib.h` are

- `void *realloc(void *p, size_t size)`, which **changes the size** of the object pointed to by `p` to a new given size, `size`, without changing its content (up to `size`), and
- `void free(void *p)`, which **de-allocates** the *heap space* pointed to by `p`.

The argument of `free` **must** be a pointer to *previously allocated dynamic memory*. For example, the following statements cause an *execution error*.

```
int i = 1;  
free(i);
```



# Free the memory

It is good practice to *free any dynamically allocated memory that is no longer needed* because

- memory usage keeps **growing** with every new allocation and
- it may help avoid memory **leaks**, i.e. allocated memory slots that are no longer referenced by a pointer.

**Suggestion:** Always check if your program leaks memory by running

```
valgrind ./a.out
```

## Example

This program produces a memory leak of 40 bytes because the corresponding pointer is *moved elsewhere*.

```
#include <stdio.h>
#include <stdlib.h>
int main() {
    int i = 1;
    int *a = malloc(10 * sizeof(int));
    for (int j=0; j<10; j++)
        *(a+j)=j;
    a = &i;
    printf("a[9]=%d\n", *(a+9));
    //free(a);
}
```

## Example

Why does running the program (without the final call to `free`) produce a *strange* output?

```
./a.out  
a[9]=-1701673536
```

If you *un-comment* the last line makes the program crash and you get

```
./a.out  
a[9]=1133786560  
free(): invalid pointer  
Aborted (core dumped)
```

# valgrind

Running the program with valgrind may help you spot the issue as you get

```
valgrind ./a.out
==1570619== Memcheck, a memory error detector
...
==1570619== Command: ./a.out
a[9]=0
==1570619== HEAP SUMMARY:
==1570619==      in use at exit: 40 bytes in 1 blocks
==1570619==    total heap usage: 2 allocs, 1 frees, 1,064 bytes allocated
==1570619== LEAK SUMMARY:
==1570619==      definitely lost: 40 bytes in 1 blocks
==1570619==      indirectly lost: 0 bytes in 0 blocks
==1570619==      possibly lost: 0 bytes in 0 blocks
==1570619==      still reachable: 0 bytes in 0 blocks
==1570619==      suppressed: 0 bytes in 0 blocks
...
==1570619== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

## references

Read more about the `malloc` family on

- [this page](#) of the online C manual,
- Section 8.7 of [The C Programming Language](#), or
- Section 9.9 of [Computer Systems: A Programmer's Perspective C](#).

# CS2850 Operating System Lab

## Week 7: Linked Lists

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

Structures

Linked lists

Implementing linked lists in C

Doubly-linked and circular lists

# Structures

Structures help *organize heterogeneous* data.

A structure is a *collection* of variables of *possibly different types*.  
The variables are called **structure members**.

A structure has a single *type-name*, called the **structure tag**.



# Defining a struct

Structure are **defined** *outside* main by specifying the **type** of each struct member, e.g.

```
struct structureTag{  
    int member1Name;  
    char *member2Name;  
};
```

Objects of type `struct structureTag` can be **declared** *inside* main as usual, e.g.

```
int main() {  
    int i;  
    struct structureTag myStruct;  
    ...  
}
```

# Accessing the struct members

There are *two equivalent ways* of accessing struct members, e.g.

```
myStruct.member1Name = 1;  
*(myStruct.member2Name) = 'h';
```

or

```
(&myStruct)->member1Name = 1  
*((&myStruct)->member2Name) = 'h'
```

**Note:** the *arrow notation* applies to the structure **address**.

## Global operations

You can **copy or assign** structures *as a unit*, e.g.

```
struct structureTag myStruct2 = myStruct;
```

You can get the **address** of the entire struct with **&**, e.g.

```
struct structureTag *pointerToMyStruct= &myStruct;
```

You can compute the **size** of a structure with **sizeof**, e.g.

```
int sizeOfMyStruct = sizeof(struct structureTag);
```

# Illegal operations

You can *not* **compare** two structures, e.g. you can *not* write

```
if(myStruct1 == myStruct2) {...}
```

You can *not* perform **arithmetic operations** between structures, e.g.  
you can *not* write

```
myStruct1 = myStruct1 + myStruct2
```

## Why not?

Structures are **composite objects** and handling them *as a unit* is not always allowed.

**Note:** all this *does not apply* to **single members**, e.g. you can have

```
if (myStruct1.member2Name == (&myStruct2)->member2Name) {  
    myStruct1.member1Name = myStruct2.member1Name + 1;  
    ((&myStruct2)->member1)++;  
}
```

# Example (1)

```
#include <stdio.h>
#define MAX 100
struct word{
    int length;
    char s[MAX];
};
int main() {
    struct word helloWorld;
    char *s = "hello, world!";
    helloWorld.length = 0;
    while (*(s + helloWorld.length) != '\0') {
        *(helloWorld.s + helloWorld.length) = *(s + helloWorld.length);
        (&helloWorld)->length++;
    }
    *(helloWorld.s + helloWorld.length) = '\0';
    struct word myStruct = helloWorld;
    printf("myStruct.s=%s\n", myStruct.s);
    printf("myStruct.length=%d\n", myStruct.length);
    printf("sizeof(myStruct)=%lu\n", sizeof(myStruct));
}
```

## Example (2)

The program above defines, initializes, copies and prints the content of a struct.

The output is

```
myStruct.s=hello, world!  
myStruct.length=13  
sizeof(myStruct)=104
```

Why do you need to initialize `myStruct.s` “*element by element*”?

# Size of a struct

The **size** in bytes of structure may depend on the order you declare the struct members, e.g. in

```
#include <stdio.h>
#define MAX 4
struct word1{
    char c;
    int v[MAX];
    char c2;
};
struct word2{
    int v[MAX];
    char c;
    char c2;
};
int main() {
    struct word1 w1;
    struct word2 w2;
    printf("sizeof(w1)=%lu\n", sizeof(w1));
    printf("sizeof(w2)=%lu\n", sizeof(w2));
}
```



# Linked lists of structures

Linked lists of structures are also called **self-referential structures**.


Each **node** in the list is a *pre-defined* object of type struct with *at least* two members:

1. the *value* of the node and
2. a *reference* to the **next** node in the list.

# Implementation

The *reference to the next node* is a **pointer** to a struct object of the *same type*, e.g.

```
struct node {  
    int val;  
    struct node *next;  
};
```



The **last element** in the list is a null pointer, NULL, so you can know when you have reached the end of the list.

# Dynamical lists

Linked lists are useful if you *do not know* the number of nodes in advance.

The idea is to **create new nodes** when new data points *arrive*.

Creating a new node requires allocating a **new memory slot** of size `sizeof(struct node)`.

The total amount of memory needed is *unknown* at compile time, i.e. you need to allocate nodes **dynamically** (using `malloc`).

## The first node

*Dynamic* linked lists are built iteratively by **linking** each new node to the *head of the list*.

A new node is declared and initialized as

```
struct node *pNode = malloc(sizeof(node));  
pNode->val = someValue;  
pNode->next = NULL;
```

# The head

A pointer, head, is needed to keep track of *where is the list head*.

Declare it as

```
struct node *head = NULL;
```

and use it to **store** the address of the list head.

When the list consists of a single node you set

```
head = pNode
```

## New nodes

As a new value, e.g. `newValue = 10`, arrives you need to

- **allocate** memory for a new node

```
struct node *cur = malloc(sizeof(struct node));
```

- **store** the new value in the node

```
cur->val = newValue;
```

- **link** the new node to the head and **move** the head to the new node

```
cur->next = head;  
head = cur;
```

## Creating the list

Initialize a list for storing a *variable number* of integers (and a 0 in the first node).

```
//headers
struct node{
    int val;
    struct node *next;
};
int main() {
    int n = getchar() - '0';
    struct node *head = NULL;
    struct node *cur = malloc(sizeof(struct node));
    cur->val = 0;
    cur->next = head;
    head = cur;
}
```

# Iterations

**Iteratively** create  $n$  new nodes and store the first  $n$  integers into them.

```
for (int i=0; i < n; i++) {  
    struct node *cur = malloc(sizeof(struct node));  
    cur->val = (i+1);  
    cur->next = head;  
    head = cur;  
}
```



## Printing the list

Print information about the nodes (from the last node, i.e. the head, to the first).

```
struct node *cur = head;
for (int i=0; i < n; i++) {
    printf("address node %d = %p\n", n-i, (void *) cur);
    printf("value node %d = %d\n", n-i, cur->val);
    printf("reference node %d = %p\n", n-i, (void *) cur
           ->next);
    printf("-----\n");
    if (cur) cur=cur->next;
}
```

## Freeing the list

*Free* the list nodes (from the last node, i.e. the head to the first).

```
for (int i=0; i < n; i++) {  
    struct node *cur = head;  
    head = cur->next;  
    free(cur);  
}
```

Freeing the nodes is important to avoid **memory leaks**.

## Other linked lists

A **doubly-linked** list is a list where each node has two references, one to the *previous* and one to the *next* node.

See [C example 7.3](#) for an example of a doubly-linked list.

A **circular** list is a *simply-linked* list where the *last* node is connected to the *first*.

See [C example 7.4](#) for an example of a circular list.

## References

More about structures can be found in the online C manual or Chapter 6 [The C Programming Language](#).

See also the code in [C example 7.1](#) on Moodle.

More about linked lists can be found in Section 6.5 of [The C Programming Language](#) or Section 10.2 of Cormen et al.'s [Introduction to Algorithms](#).

See also the code in [C example 7.2](#), [C example 7.3](#), and [C example 7.4](#) on Moodle.

# CS2850 Operating System Lab

## Week 8: Pipes

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

Inter-process communication

Anonymous pipes

C implementation

# Processes

A *process* is the operating system's **abstraction** of a running program.

When multiple processes run **concurrently**, each process appears to have *exclusive* use of the hardware.

The process **context** is the *state information* the process needs to run.

The OS keeps the context associated with concurrent processes **separate**.

# Inter-process communication

Can two processes **communicate**?

To connect the context of two processes you need an IPC *system call*.

**Example:** IPC occurs when a parent process *waits* for a child to terminate (by calling `wait`).

The child sends a *termination signal* to the parent that *can be decoded* through the macro `WEXITSTATUS`



## example

IPC between a child and its parent.

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>
#include <stdlib.h>
int main() {
    int status, pid;
    int N = getchar() - '0';
    int *vOut = malloc(sizeof(int) * N);
    int *vPid = malloc(sizeof(int) * N);
    pid_t pid;
    for (int j = 0; j < N; j++) {
        if (! (pid=fork())) return rand() % (1 + j);
        *(vPid + j) = wait(&status);
        *(vOut + j) = WEXITSTATUS(status);
    }
    for (int k = 0; k < N; k++) printf("%d-pid=%d returns %d \n",
        k, *(vPid + k), *(vOut+k));
    free(vOut);
    free(vPid);
}
```

# More general IPC

There exist different types of IPC:

- *Signals*: **simple** messages that are not used to transfer data,
- *Anonymous pipes*: **unidirectional** data channels between *related processes*,
- *Named pipes*: more flexible pipes that are treated like a **file**,
- *Shared memory*: a memory block that can be **accessed** by multiple concurrent processes,
- ...

See [Section 6 of the Linux online manual](#) or [Wikipedia page on IPC](#) for more details about UNIX and general IPC.

# Linux channels

A channel is an IPC model for **message passing**.

It has a **shareable** reference that allows more processes to access it.

Processes can access the channel to **write** a message *in* it or **read** a message *from* it.

Channels normally have **two distinct references**, one for *reading* or *writing*.

They behave like **one-way** tunnels with two *gates*: an *entry* on one side and an *exit* on the other.

# Writing and reading

Processes connected through a channel know the reference of (at least) one end of the channel so that

- one of them can *push* data into the channel and
- one of them can *pull* data out of the channel

i.e. reading and writing happen at different ends.

Written data *travel* through the tunnel before they are read and arrive in the **order** they are sent (FIFO).

Data written into the channel are *buffered* **by the OS** until they are read from the other end.

# Implementation

In C, you can implement 3 types of channels:

- *half-duplex UNIX pipes*: communication between related processes, e.g. a parent and a child or the children of the same parent,
- *FIFOs named pipes*: communication between two independent processes,
- *sockets*: communication between different computers,
- ...

## UNIX half-duplex (*anonymous*) pipes

Anonymous pipes are the **eldest** of the IPC tools.

A **half-duplex** pipe connects processes that *share a related ancestor*, i.e.

- a parent and child or
- children of the same parent.

Defining pipes requires **kernel-level operations**, i.e. a pipe is created by making a *system call*, e.g. by calling `pipe()` in a C program.

**2-way** pipes can be created by opening two pipes.

# Example

Pipes are used in the **UNIX shell** to connect the *output* of one process to the *input* of another one, e.g. in

```
ls | sort
```

Intuitively, you can imagine data flow through the channel from the *left* to the *right* both the `ls` and `sort` processes are *children of the same process*, i.e. the shell.


## Creating a pipe: step 1

The system call `pipe()` creates two **file descriptors** associated with the *reading* and *writing* ends:

- `fd0` for *reading*,
- `fd1` for *writing*.

The new pipe initially connects a process to **itself**.

Data *written to* and *read from* the pipe travels through the **kernel**.





## Step 1: code

```
#include <stdio.h>
#include <unistd.h>
#define MAXCHARS 100
int main() {
    int fd[2];
    pipe(fd);
    printf("fd[0]=%d, fd[1]=%d\n", fd[0], fd[1]);
    char buf[MAXCHARS];
    *buf = '\0';
    printf("buf = %s\n", buf);
    int nbytes = write(fd[1], "hello world", 12);//write to fd[1]
    printf("%d written to fd[1]\n", nbytes);
    nbytes = read(fd[0], buf, nbytes);
    printf("%d read from fd[0]\n", nbytes);// and read from fd[0]
    printf("buf = %s\n", buf);
    return 0;
}
```

## Step 1: output

The output of the program is

```
fd[0]=3, fd[1]=4  
buf =  
12 bytes written to fd[1]  
12 bytes read from fd[0]  
buf = hello world
```

# File descriptors

**File descriptors**, e.g. `fd[0]` and `fd[1]`, are used by *low-level* I/O functions as `read` and `write`.<sup>1</sup>

The *standard library* functions refer to files through file *handles* or **streams**, i.e. *structures* that contain more information about the file.

Given a file descriptor, you can obtain the corresponding stream and *vice versa* through.

```
FILE *streamPointer = fdopen (f0, openMode)
int f1 = fileno(streamPointer)
```

where `f0` and `f1` are (equal) integers, `openMode` a string, e.g. `"w"` or `"r"`, and `streamPointer` a pointer to a `FILE` structure.

---


<sup>1</sup>See [Slides 4](#) or [Section 13.4 of the GNU online manual](#) for more details about low-level I/O

## Creating a pipe: step 2

What happens if the process generates a **child** process *after creating* the pipe?

As `fork` makes an *exact copy* of the parent process, the pipe **file descriptors** are copied from the parent into the child,

Both processes have then **access** to the pipe and can use it to *communicate*.



## Step 2: code

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int copyString(char *out, char *in);
int main() {
    int fd[2], nBytes;
    pipe(fd);
    char out[100], in[100], author[100];
    if (!fork()) {//child process
        sleep(1);
        nBytes=copyString(in, "world");
        copyString(author, "child");
    }
    else { //parent process
        nBytes=copyString(in, "hello");
        copyString(author, "parent");
    }
    nBytes = write(fd[1], in, nBytes);//write to fd[1]
    printf("the %s writes to fd[1]: %s (%d)\n", author, in, nBytes);
    wait(NULL);
    nBytes = read(fd[0], out, nBytes);
    printf("the %s reads from fd[0]: %s (%d)\n",author, out, nBytes);
}
```

## Step 2: copyString and program output

copyString is defined by

```
int copyString(char *out, char *in) {  
    int k = 0;  
    while (*(in + k) != '\0') {  
        *(out + k) = *(in + k);  
        k++;  
    }  
    *(out + k) = '\0';  
    return k;  
}
```

The program produces the following output

```
the parent writes to fd[1]: hello (5)  
the child writes to fd[1]: world (5)  
the child reads from fd[0]: hello (5)  
the parent reads from fd[0]: world (5)
```

# Problems

Why do we get such an output?

The child process inherits both **file descriptors**, `fd[0]` and `fd[1]` both processes can write to **and** read from the channel.

While the child is sleeping the parent **writes** to the channel.

The **order** of the messages is kept in the channel (channels are **FIFO** containers) so messages written by the parent and the child may get **mixed**.


## Creating a pipe: step 3

You must **decide** *who writes* and *who reads* i.e. in which direction data will travel.

The processes should *close* the **unused end** of the pipe, i.e.

- close **fd[0]** if the process will be *writing* and
- close **fd[1]** if the process will be *reading*

**Example:** The child *writes* and the parent *reads*.





# Pseudo-code

`createPipe.c` creates a message-passing system between a parent process and its child in a few steps.

- The parent creates a **half-duplex** pipe. (3)
- The parent generates a **child** process by calling `fork`. (5)
- Both processes **close** the *unused* ends of the pipe (6 and 15).
- The child's **sends** messages to the parent by calling `writeMessages`. (9)
- The parent **starts reading** the child messages and prints them on `stdout` by calling `readMessages`. (19)
- The child sends a **signal**, e.g. a message containing the *keyword* `quit`, to tell the parent that the *last message* has been sent. (10)
- The child and the parents **close** the channel. (11 and 20 )

## createPipe.c (main)

```
int main() { 1
    int fd[2]; 2
    pipe(fd); 3
    char process[MAXCHARS]; 4
    if (!fork()) { 5
        close(fd[0]); 6
        copyString(process, "the child"); 7
        writeMessage(fd[1], "hello", process); 8
        writeMessage(fd[1], "world", process); 9
        writeMessage(fd[1], "stop", process); 10
        close(fd[1]); 11
        return 0; 12
    } 13
    else { 14
        close(fd[1]); 15
        copyString(process, "the parent"); 16
        int end = 0; 17
        while(end == 0) 18
            end = readMessage(fd[0], process); 19
        close(fd[0]); 20
    } 21
} 22
```

## createPipe.c (Headers and subroutines declaration)

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#define MAXCHARS 100
int readMessage(int fd, char *reader);
int writeMessage(int fd, char *buf, char *author);
int compareString(char *s1, char *s2);
int copyString(char *in, char *out);
int stringLength(char *s);
```

## createPipe.c (Read and write functions)

```
int writeMessage(int fd, char *buf, char *author) {
    int nbytes = strlen(buf);
    printf("%s writes to fd[1]: %s (%d) \n", author, buf, nbytes);
    nbytes = write(fd, buf, nbytes + 1);
    return 0;
}

int readMessage(int fd, char *reader) {
    char c = ' ';
    int i = 0;
    char buf[MAXCHARS];
    while (c != '\0') {
        read(fd, &c, 1);
        buf[i++] = c;
    }
    buf[i] = '\0';
    if (strcmp(buf, "stop") == 0) return -1;
    printf("%s read from fd[0]: %s (%d) \n", reader, buf, strlen(buf));
    return 0;
}
```

## Step 3: code (string functions)

```
int copyString(char *out, char *in) {
    int k = 0;
    while (*(in + k) != '\0') {
        *(out + k) = *(in + k);
        k++;
    }
    *(out + k) = '\0';
    return k;
}

int stringLength(char *s) {
    int i = 0;
    while (*(s + i) != '\0')
        i++;
    return i;
}

int compareString(char *s1, char *s2) {
    int i = 0;
    while (*(s1 + i) != '\0' && *(s2 + i) != '\0')
        if (*(s1 + i) != *(s2 + i)) return -1;
        else i++;
    if (*(s1 + i) == *(s2 + i)) return 0;
    else return -1;
}
```

### Step 3: createPipe.c (output)

```
the child writes to fd[1]: hello (5)
the child writes to fd[1]: world (5)
the child writes to fd[1]: stop (4)
the parent read from fd[0]: hello (5)
the parent read from fd[0]: world (5)
```

**Note:** The implementation of `strlen` makes the program print the actual size of a word, excluding the null-termination character.

Would it be possible for the parent to read *hello* *before the child writes world*?

# CS2850 Operating System Lab

## Week 9: Threads

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# Outline

Threads

POSIX threads

Mutexes



# Threads

Computers are **busy** machines.

Threads help them perform **several tasks** at the same time.

A single **process** can solve *different related problems* by starting multiple threads.

Threads **share** the same *global* memory, e.g. data, but have their own **stack** (*automatic variables*).

# Address space

A process memory contains

- **process-wide** resources, e.g. the program instructions and global data, and
- **execution state** information, e.g. the program counters and stack.

All threads can access process-wide resources but have **private** execution state information.

# POSIX threads

There are **differences** between thread packages. Here we focus on **Pthreads** (for *POSIX* threads).

The **Portable Operating System Interface (POSIX)** is a set of *types and functions* that you can call from a C program (under UNIX).

Pthreads let you *divide* a program in *sub-tasks* and execute them sequentially or **in parallel**.

# Pthreads in a C program

How to *create* multiple threads in C?

In a C program, you create new threads with a POSIX function called `pthread_create`.

You also need to

- define the `task` assigned to each thread (normally a subroutine of your program) and
- call a function that `merges` the output of threads when they return.

# All threads are equal

There is *no hierarchy* between running threads, e.g. no parent-child distinctions.

Each thread executes **independently**: the actual execution order is *unpredictable*.

You **synchronize** them *explicitly* by **waiting** until they return.

Handling threads with POSIX *functions* requires to

- *include* the header **pthread.h** and
- *compile* your program with a **-pthread** flag.

# Creating a Pthread

You create a new thread with<sup>1</sup>

```
int pthread_create(pthread_t *thread, const
    pthread_attr_t *attr, void *subroutine, void *arg)
```

- thread points to a *buffer* that stores the thread **identifier**,
- attr points to a struct that specifies various **features** of the new thread (write NULL for the default),
- subroutine points to the **subroutine** executed by the new thread, and
- arg points to the **parameter** of start\_routine.

---

<sup>1</sup>See more about pthread\_create on [this page](#) of the [Linux online manual on Pthreads](#).

# Terminating a Pthread

You **wait** for a thread to terminate and see its **return value** with

```
int pthread_join(pthread_t thread, void **retval)
```

- thread is the thread **identifier** and
- retval is the location where the **exit status** of the target thread is copied.

A thread **terminates** when it reaches a return statement or calls

```
void pthread_exit(void *retval)
```

# Synchronization

You can ensure that one *event* (in one thread) happens **before** another event (in a concurrent thread) by **synchronizing** the two threads.

For example, you can **wait** for one thread to *return* by calling `pthread_join` in the main program.

See more about thread synchronization in Chapter 3 of *PThreads Programming* by Farrel et al..




# Threads or processes?

In the following examples, we implement the same program using

- no threads,
- a single thread,
- two threads, and
- two child processes.

# A process with no threads



- Text contains the **program instructions**.
- Data are the **global data** needed to run the program.
- Heap is for **dynamic** memory allocation.
- Stack is for storing **automatic variables**.
- Registers, Identity, and Resources contain the **info** needed to execute.


## Example: main (simple version)

```
#include <stdio.h>
void sayGoodbye(int n1, int n2);
int brian(int n), dennis(int n);
int main(){
    int nBrian = 3, nDennis = 2;
    while(nBrian + nDennis > 0) {
        if (nBrian) nBrian = brian(nBrian);
        if (nDennis) nDennis = dennis(nDennis);
        sayGoodbye(nBrian, nDennis);
    }
}
int brian(int n) {
    for (int i = 0; i < n; i++) printf("%d-hi Dennis!\n", i + 1);
    return n - 1;
}
int dennis(int n) {
    for (int i = 0; i < n; i++) printf("%d-hi Brian!\n", i + 1);
    return n - 1;
}
void sayGoodbye(int n1, int n2) {
    printf("%d-goodbye!\n", n1 + n2);
}
```

# Output

```
./a.out  
1-hi Dennis!  
2-hi Dennis!  
3-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
3-goodbye!  
1-hi Dennis!  
2-hi Dennis!  
1-hi Brian!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```

# A process with a single Pthreads



Stack and Registers are now part of the thread.

Each *active thread* has a **reserved stack frame** and **specific machine registers**.

## Example: main (Pthread compatible)

pthread\_create takes three free arguments

```
int pthread_create(pthread_t *tName, NULL, void *f(void
    *), void *arg)
```

void \*f(void \*) is (a pointer to) a function of type void which takes a pointer to void as a single argument.

The function arguments should be passed through a single pointer to void.

You need to rewrite brian and dennis as

```
void *brian(void *n)
void *dennis(void *n)
```

## New version of brian and dennis

To meet the requirement of `pthread_create` you need a structure holding the original integer argument,

```
struct arg {  
    int n;  
};
```

In the first line of `f`, cast the void pointer to be a pointer to a variable of type `struct arg`<sup>2</sup>

```
struct arg *p = n;
```

In main, `n` becomes `(&argName)->n`, where `argName` is an object of type `struct arg`.

---

<sup>2</sup>To access the original arguments with `p->n`.

# New version of brian and dennis

You obtain

```
void *brian(void *nIn) {  
    struct arg *nStack = nIn;  
    for (int i=0; i<nStack->n; i++)  
        printf("%d-hi Dennis!\n", i + 1);  
    nStack->n = nStack->n - 1;  
    return NULL;  
}  
  
void *dennis(void *nIn) {  
    struct arg *nStack = nIn;  
    for (int i=0; i<nStack->n; i++)  
        printf("%d-hi Brian!\n", i + 1);  
    nStack->n = nStack->n - 1;  
    return NULL;  
}
```



# New version of brian and dennis

The new version of main is

```
#include <stdio.h>

struct arg{
    int n;
};

void *brian(void *n), *dennis(void *n), sayGoodbye(int n1, int n2);

int main(){
    struct arg nBrian = {3}, nDennis = {2};
    while(nBrian.n + nDennis.n > 0) {
        if (nBrian.n) brian(&nBrian);
        if (nDennis.n) dennis(&nDennis);
        sayGoodbye(nBrian.n, nDennis.n);
    }
}
```

sayGoodbye can be left unchanged.

## Example (1 running Pthread)

```
#include <stdio.h>
#include <pthread.h>

struct arg{
    int n;
};

void *brian(void *n), *dennis(void *n), sayGoodbye(int n1, int n2);

int main(){
    pthread_t t1;
    struct arg nBrian = {3}, nDennis = {2};
    while(nBrian.n + nDennis.n > 0) {
        if (nBrian.n) pthread_create(&t1, NULL, brian, &nBrian);
        pthread_join(t1, NULL);
        if (nDennis.n) pthread_create(&t1, NULL, dennis, &nDennis);
        pthread_join(t1, NULL);
        sayGoodbye(nBrian.n, nDennis.n);
    }
}
```


All auxiliary functions are defined as above.

# Output (1 running Pthread)

The execution is as before

```
./a.out  
1-hi Dennis!  
2-hi Dennis!  
3-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
3-goodbye!  
1-hi Dennis!  
2-hi Dennis!  
1-hi Brian!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```

# A process with two threads



Both threads have *their private copy* of the machine **registers** and can use variables or file descriptors in the *process-wide areas*.

The two threads execute at **different locations**.

## Example (2 running Pthreads)

```
#include <stdio.h>
#include <pthread.h>

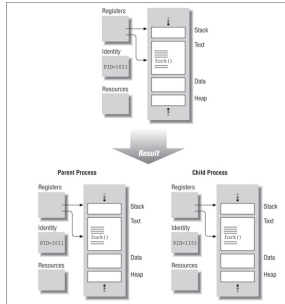
struct arg{
    int n;
};

void *brian(void *n), *dennis(void *n), sayGoodbye(int n1, int n2);

int main(){
    struct arg nBrian = {3}, nDennis = {2};
    pthread_t t1, t2;
    while(nBrian.n + nDennis.n > 0) {
        if (nBrian.n) pthread_create(&t1, NULL, brian, &nBrian);
        if (nDennis.n) pthread_create(&t2, NULL, dennis, &nDennis);
        pthread_join(t1, NULL);
        pthread_join(t2, NULL);
        sayGoodbye(nBrian.n, nDennis.n);
    }
}
```

As t1 starts before t2 and the task performed by brian is *computationally cheap*, the output on stdout is unchanged.

# Two processes vs. two threads



The address space of two concurrent processes created by `fork` is *completely private* (no shared areas).

The two processes can interact via a **channel** (anonymous pipe).

## Example: main (2 running child processes, slowBrian)

```
int main() {  
    int status;  
    while(nBrian.n + nDennis.n > 0) {  
        int pidBrian = fork();  
        if (!pidBrian) {  
            if (nBrian.n) brianSlow(&nBrian);  
            return nBrian.n;  
        }  
        int pidDennis = fork();  
        if (!pidDennis) {  
            if (nDennis.n) dennis(&nDennis);  
            return nDennis.n;  
        }  
        waitpid(pidBrian, &status, 0);  
        nBrian.n = WEXITSTATUS(status);  
        waitpid(pidDennis, &status, 0);  
        nDennis.n = WEXITSTATUS(status);  
        sayGoodbye(nBrian.n, nDennis.n);  
    }  
}
```

Include `unistd.h` and `sys/wait.h` to use `fork` and `wait`.

## Output (2 running child processes, slow Brian)

The output is the same as before

```
./a.out  
1-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
2-hi Dennis!  
3-hi Dennis!  
3-goodbye!  
1-hi Dennis!  
1-hi Brian!  
2-hi Dennis!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```



## Another example of synchronization: Mutexes

Variables of type `pthread_mutex_t` allow you to synchronize the **access** to some *area of your code*.

The protected area is **mutually exclusive**. While one thread is in the protected area, other threads cannot

- *execute* the protected code and
- *lock or unlock* the Mutex.

# Mutex-protected variables

To use a Mutex to regulate the *updates of shared variable*, *x*,

1. create and initialize a **mutex variable**, *m*, using

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;  
(m should be declared outside main).
```

2. in the thread subroutine, *lock and unlock* the mutex before updating *x*, e.g. write

```
pthread_mutex_lock(&m)  
x= ...  
pthread_mutex_unlock(&m)
```

## Example: (2 running threads with thread-safe subroutines)

The *thread-safe* version of slowBrian and dennis is

```
void *brianSlow(void *n) {
    struct arg *nl = n;
    int t = 0;
    pthread_mutex_lock(&m); //start of the mutual exclusive area
    for (int i=0; i<(nl->n); i++) {
        printf("%d-hi Dennis!\n", i + 1);
        while (t < 100000) t++;
    }
    pthread_mutex_unlock(&m); //end of the mutual exclusive area
    nl->n = nl->n - 1;
    return NULL;
}

void *dennis(void *n) {
    struct arg *nl = n;
    pthread_mutex_lock(&m); //start of the mutual exclusive area
    for (int i=0; i<(nl->n); i++) printf("%d-hi Brian!\n", i + 1);
    pthread_mutex_unlock(&m); //end of the mutual exclusive area
    (nl->n) = (nl->n) - 1;
    return NULL;
}
```

## Output (2 running threads, thread-safe subroutines)

```
./a.out  
1-hi Dennis!  
2-hi Dennis!  
3-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
3-goodbye!  
1-hi Dennis!  
2-hi Dennis!  
1-hi Brian!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```

# CS2850 Operating System Lab

## Week 10: Review of pointers, strings, and process control

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

# outline

Pointers

Strings

Pointers and functions

fork

# Pointers (1)

Pointers are **variables** that store the memory address of other variables.

*Address operator:* `&i` is the **address** of `i`.

*Dereferencing operator:* `*ip` is the **value** of the variable stored at the address `ip`.

You can **print** on stdout the *unpredictable value* of `ip`, i.e. the address stored in `ip`, with

```
printf("%p", ( void *) ip)
```

## Pointers (2)

The **type** of a pointer is the type of variable stored at the pointed address.

The data type of a variable specifies how to **interpret** the data stored in the variable.

A pointer of a given type **cannot** store the address of variables of a different type, e.g.

```
int i = 1;  
char *pc = &i
```

is **illegal** and produces compilation errors.



## Example: int, char, int\*, and char\* (1)

```
#include <stdio.h>
int main() {
    int i = 10;
    int *pi = &i;
    char c = 'a';
    char *pc = &c;
    printf("sizeof(pi)=%lu\n", sizeof(pi));
    printf("sizeof(pc)=%lu\n", sizeof(pc));
    printf("pi=%p and  &i = %p\n", (void *) pi, (void *) &
        i);
    printf("pc=%p and &c = %p\n", (void *) pc, (void *) &c
        );
    printf("*pi=%d and i = %d\n", *pi, i);
    printf("*pc=%c and c = %d\n", *pc, c);
}
```

# Output

```
./a.out
sizeof(pi)=8
sizeof(pc)=8
pi=0x7ffd17e6b974 and &i = 0x7ffd17e6b974
pc=0x7ffd17e6b973 and &c = 0x7ffd17e6b973
*pi=10 and i = 10
*pc=a and c = 97
```

# Strings (1)

**String constants** are *null-terminated* arrays of char, e.g.

```
char *s = "hello world"
```

s is a pointer to char that stores the address of the first element of the string "hello world".

The last char of s is the null character '**\0**'.

## Strings (2)

String *constants* and character arrays, e.g. `s` and `as` defined by

```
char *s = "hello world";  
char as[12] = "hello world";
```

are *not equivalent*.

`s` is a pointer (a variable) and can be *assigned to a different address*.

`as` is an array and always refers to *fixed allocated* storage.

## Strings (3)

`s[7]` and `as[7]` both return `'w'` but

```
s[7] = 'x';
```

causes a Segmentation Fault, while

```
as[7] = 'x';
```

is **legal** and sets the 8th character of `as` to `'x'`, i.e. `'o'` becomes `'x'`.

## Strings (4)

You can use `printf("%s", s)` and `printf("%s", as)` to **print** the strings `s` and `as` **up to** their null-termination character.

In C, there are *no built-in operators* for processing an **entire** string as a unit, e.g.

```
char as[12];  
as = 'hi world';
```

is *illegal*, but you can *re-assign* `s` to point to another string constant, e.g.

```
char *s;  
s = 'hi world';
```

## Strings (6)

You can print **a portion** of `s` or `as` by specifying the address of a single character, e.g.

```
s = "hi world";  
printf("(s + 3) = %s", (s + 3));
```

prints (s + 3) = world.

You can **cut** a character array by setting one element to `'\0'`, e.g.

```
char as[10] = "hi world";  
*(as + 2) = '\0';  
printf("as = %s and (as + 3) = %s\n", s, (s + 3));
```

prints as = hi and (as + 3) = world.

## Example: char\* and char[13]

```
#include <stdio.h>
int main() {
    char *s = "hello world";
    char sa[13] = "hello world";
    //s[7] = 'x'; is illegal
    sa[7] = 'x'; //is legal
    printf("s[7] = %c\n", s[7]);
    printf("sa[7] = %c\n", sa[7]);
    printf("(s + 1) = %s\n", s + 1) ;
    printf("(sa + 1) = %s\n", sa + 1);
    s = "hi world";
    printf("s = %s\n", s);
    //sa = s; is illegal
    s = &sa[6]; //is legal and equivalent to s = sa + 6;
    printf("s = &sa[6] => s=%s\n", s);
    printf("(sa + 6) = %p and s = %p\n", (void *) (sa + 6),
        (void *) s);
}
```



# Output

```
./a.out  
s[7] = o  
sa[7] = x  
(s + 1) = ello world  
(sa + 1) = ello wxrld  
s = hi world  
s = &sa[6] => s=wxrld  
(sa + 6) = 0x7ffd23e608c1 and s = 0x7ffd23e608c1
```

## Example: pointer arithmetic

```
#include <stdio.h>
int main() {
    char as[100] = "one two three four five";
    char vs[100] = "one two three four five";
    char *s = "one two three four five";
    for (int i = 0; i < 13; i++) s++;
    as[13] = '\\0';
    printf("as = %s\\n", as);
    printf("vs = %s\\n", vs);
    printf("s = %s\\n", s);
}
```

### Output:

```
as = one two three
vs = one two three four five
s = four five
```

# Pointers and functions (1)

The de-referencing operator can be used in **function declarations**, e.g.

```
double f(char *c);
```

says that f **returns** a double and the **argument** of f is a pointer to char.

Arguments are passed to functions **by value**. There is *no direct way* for the called function to **modify** a variable in the calling function.

To change the value of a variable in the calling function, pass **the address** of that variable.

## Pointers and functions (2)

**Arrays** are represented by *the location of the initial element*.

When you pass an array to a function, you will see the changes in `main`, e.g.

```
#include <stdio.h>
void initialise(int *a) {
    int i;
    for (i = 0; i < 10; i++) *(a + i) = i;
}
int main() {
    int i, a[10];
    initialise(a);
    for (i = 0; i < 10; i++) printf("%d ", *(a + i));
}
```

prints 0 1 2 3 4 5 6 7 8 9.

## Example: swapping int

```
#include <stdio.h>
void swap(int *pi, int *pj);
void swapWrong(int i, int j);
int main() {
    int x = 1, y = 2, *px = &x, *py = &y;
    printf("x = %d, y = %d, px = %p, and py = %p\n", x, y, (void *) px, (void *) py);
    printf("calling swapWrong ... \n");
    swapWrong(x, y);
    printf("x = %d, y = %d, px = %p, and py = %p\n", x, y, (void *) px, (void *) py);
    printf("calling swap ... \n");
    swap(&x, &y);
    printf("x = %d, y = %d, px = %p, and py = %p\n", x, y, (void *) px, (void *) py);
}
```

## Output:

```
./a.out
x = 1, y = 2, px = 0x7fff73496720, and py = 0x7fff73496724
calling swapWrong ...
x = 1, y = 2, px = 0x7fff73496720, and py = 0x7fff73496724
calling swap ...
x = 2, y = 1, px = 0x7fff73496720, and py = 0x7fff73496724
```

## Auxiliary functions

A function to swap the *content of two addresses*.

```
void swap(int *pi, int *pj) {  
    int temp = *pi;  
    *pi = *pj;  
    *pj = temp;  
}
```

A function to swap two *stack variables*.<sup>1</sup>

```
void swapWrong(int i, int j) {  
    int temp = i;  
    i = j;  
    j = temp;  
}
```

---

<sup>1</sup>All changes are discarded when the function returns.

# Process control (1)

`unistd.h` and `sys/wait.h` contains the functions you need to **control** a process from a C program.

`int getpid()` returns the **process identifier** (PID) of the running process.

`int fork()` creates a **child** process that is a copy of the running process.

`void exit()` or return **terminate** the current process.

`int wait(int * status):` to **wait** for the *first* child that terminates (call it several times if you have generated more than one child). <sup>2</sup>

---

<sup>2</sup>See more about process control in [Section 26 of the GNU online manual](#).

## Process control (2)

The OS assigns PIDs in **increasing order** but their values are unpredictable.

In C, `fork` is *called once but returns twice*.

The return value is  $-1$  if the process creation failed,  $0$  in the child process, and the child's PID in the parent process.



## Process control (3)

When `fork` returns, the child address space is a copy of the parent address space, but the parent and the child have **different PIDs**.

The **return values** of `fork` allow you to make conditional statements.

The address spaces of the child and the parent are *private*, i.e. the parent cannot see the **changes** made by the child, and vice versa.

The parent and the child **run concurrently**. The execution order is unpredictable.

## Example: private but identical address spaces

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main() {
    int j = 0;
    printf("PID = %d, j = %d, &j = %p \n", getpid(), j, (void *) &j);
    int pid = fork();
    if (!pid) j = 1;
    else j = 2;
    printf("PID = %d, j = %d, &j = %p \n", getpid(), j, (void *) &j);
}
```

### Output:

```
./a.out
PID = 2789388, j = 0, &j = 0x7fff5ea69010
PID = 2789388, j = 2, &j = 0x7fff5ea69010
PID = 2789389, j = 1, &j = 0x7fff5ea69010
```