

# Operating Systems Lab - Week 2: exercise

## - with answers

This lab is about C types, variables, and functions. You will see in practice how numbers and characters are represented in C and how you can define and call functions. You will also write programs that parse terminal input/output.

### Set up

We suggest you edit, save, and compile the programs you write for this lab session in `CS2850Labs/week2` a dedicated sub-folder of the directory you created for the first lab session of the term, on the teaching server, `linux.cim.rhul.ac.uk`.

On the course Moodle, page you can find more details about [connecting](#) to `linux.cim.rhul.ac.uk` and [editing and compiling](#) your code from the command-line and [debugging](#) your programs using Valgrind.

## 1 Variables

Similarly to other programming languages, you can use variables of different *types* and different *storage classes*. This [list of primitive data types](#) contains all variables you can declare and use in C. In this section, you will write a program that prints on screen the size in bytes of the most common C types, i.e.

`char, int, unsigned int, float`

### 1.1 Integers

Start by declaring and initializing a variable of type `int` and print its value as in the following program

```
#include <stdio.h>
int main() {
    int a = 10;
    printf("a=%d\n", a);
}
```

1  
2  
3  
4  
5

Write the code above into a file called, `printInt.c` and compile and run it to check that it prints

`a=10`

1

on screen. You can modify and recompile `printInt.c` as suggested in the following questions:

1. What happens if the variable is declared outside `main`?

**Answer:** The variable is automatically declared as a static variable but the output does not change.

2. What happens if you add a non-integer part in the initialisation of `a`, e.g. if you replace Line 3 with `a = 10.1234`?

**Answer:** The non-integer part is truncated.

3. What happens if you initialised `a` with a *very large* value, e.g. if you replace Line 3 with `a = 2147483647` and `a = 2147483648`?

**Answer:** The second initialization produces a compilation error/warning because the value is out of the `int` range

The following code produces the same output as the program above.

```

#include <stdio.h>
int a = 10;
void printValue() {
    printf("a=%d\n", a);
}
int main() {
    printValue();
}

```

Copy the new program into a new file, **printInt2.c**, compile it, and run it to check that its output on screen is indeed

```
a=10
```

Modify **printInt2.c** as suggested in the following questions:

1. What happens if the variable is declared *inside* **main**?

**Answer:** You get a compilation error because the function does not know anything about **a**

2. What happens if you change the value of **a** inside **main**, e.g. if you add

```
a = 11
```

just before Line 7?

**Answer:** The function see the new value because the variabvle is global

3. What happens if you change the value of **a** inside the definition of **printValue**, e.g. if you add

```
a = 11
```

just before Line 4?

**Answer:** The function prints the updated value because the variable is global

## 1.2 unsigned int, char, and float

Write a modified version of **printInt.c** called **printTypes.c**, that prints

```

au=2147483648
ac=*
af=0.123456

```

on screen and where **au** is declared as an **unsigned int**, **ac** as a **char**, and **af** as **float**. To obtain the correct output you should also use the correct format identifiers, **%u** for **unsigned int**, **%c** for **char**, and **%f** for **float** in the call of **printf**. Have a look at this list of formatting symbols for more details.

**Answer:**

```

#include <stdio.h>
int a = 1234;
unsigned int au = 2147483648;
char ac = '*';
float af = 0.123456;
void printValue() {
    printf("a=%d\n", a);
    printf("au=%u\n", au);
    printf("ac=%c\n", ac);
    printf("af=%f\n", af);
    printf("a=%e\n", (float) a);
}

```

```

}
int main() {
    printValue();
}

```

What happens if you use the `int` format, `%d`, instead of `%u` when you call `printf` in `printUnsigned.c`?

**Answer:** You get

```

au=-2147483648

```

Try also to print the value of the variables as an unsigned octal number, with `%o`, an unsigned hexadecimal number, with `%x`, and a floating-point number in exponential notation, with `%e`. Which conversions are allowed and which lead to a compilation error if the program is compiled using the `-Werror -Wall` flags?

**Answer:** A `float` cannot be printed as an octal or hexadecimal number, e.g. you get

```

format '%o' expects argument of type 'unsigned int', but argument 2 has type 'double'
format '%x' expects argument of type 'unsigned int', but argument 2 has type 'double'

```

and integers cannot be printed in the exponential notation, e.g. you get

```

format '%e' expects argument of type 'double', but argument 2 has type 'int'
format '%e' expects argument of type 'double', but argument 2 has type 'unsigned int'

```

Force the conversion by including a type *cast* in the *second argument* of `printf` as in the following example

```

#include <stdio.h>
int a = 1234;
void printValue() {
    printf("a=%e\n", (float) a);
}
int main() {
    printValue();
}

```

### 1.3 Sizes

The size of a given type can be obtained by calling the operator `sizeof(type)`, e.g.

```

unsigned long int sizeOfChar = sizeof(char);

```

idem with `int`, `unsigned int`, or `float`, or by letting the argument of `sizeof` be a pre-declared variable, e.g.

```

char a;
unsigned long int sizeOfChar = sizeof(a);

```

See Section A7.4.8 of [The C Programming Language](#) or [Section 3.11](#) of the GNU Online Manual for more details about the `sizeof` operator. Write a program, `sizeofTypes.c`, that prints on the terminal the size in bytes of a `char`, an `int`, an `unsigned int`, and a `float`. Your program should print the size of each type on a different line, with each line being of the form

```

the size of a long int is 8 bytes

```

Note that the output of `sizeof` is an `unsigned long int`.

**Answer:**

```

#include <stdio.h>
int main() {
    printf("the size of a char is %lu bytes \n", sizeof(char));
    printf("the size of an int is %lu bytes \n", sizeof(int));
    printf("the size of an unsigned int is %lu bytes \n", sizeof(unsigned int));
}

```

```
printf("the size of a float is %lu bytes \n", sizeof(float));
}
```

6  
7

What happens if you use `sizeof` to get the memory size associated with an array? Modify your program so that it prints two extra lines reporting the size in bytes of a 10-dimensional array of `char` and `int` declared as

```
int vInt[10];
char vChar[10];
```

**Answer:** The size of an `int` or a `char` is multiplied by the number of items in the array.

## 1.4 Signed or unsigned char (optional)

The conversion of characters to integers depends on whether the compiler treats the variables of type `char` as signed or unsigned quantities. Try to understand if on your system they are signed or unsigned by looking at the error messages produced by `gcc -Wall -Werror -Wpedantic` when you compile a program such as

```
#include <stdio.h>
int main() {
    char a = 150;
    unsigned char b = 150;
    printf("a=%d and b=%d\n", a, b);
}
```

1  
2  
3  
4  
5  
6

**Answer:** The compiler prints an error because 150 is out of range if `char` is a *signed variable of 1 byte*. The conversion of a variable of type `int` into type `char` may cause some information to be lost. Copy, compile, and run the following program:

```
#include <stdio.h>
int main() {
    int a = 128;
    char c;
    c = a;
    a = c;
    printf("a=%d\n", a);
}
```

1  
2  
3  
4  
5  
6  
7  
8

What do you observe? Can you explain why all problems disappear if initialise `a` with the value 127?

**Answer:** The value of `a` becomes `-127`. If `a` is initialised with 127 it keeps its value because 127 is within the range of a signed `char`.

## 2 Terminal input/output: `getchar` and `putchar`

In this section, you will write a program that transforms all lower case letters of an input string into upper case letters. The standard library contains functions for reading or writing one character at a time:

1. `getchar()`, which reads the next input character and returns it, and
2. `putchar(c)`, which prints the character `c` on the terminal.

Read, and try to guess what the following program does

```
#include <stdio.h>
int main() {
    char c;
    while ((c = getchar()) != 'q') {
```

1  
2  
3  
4

```

        putchar(c);
    }
}

```

Copy the code into a new file called `inputOutput.c`, compile, and run it to understand how it works by typing random character on the screen when the program starts.

**Answer:** A typical run produces

```

cim-ts-node-01$ ./a.out
2
2
a
a
sa
sa
e
e
v
v
askkj
askkj
sdmnsn salkjsdd
sdmnsn salkjsdd
q
cim-ts-node-01$

```

## 2.1 Change the *exit* keyword

When you run the program in `inputOutput.c`, the terminal shows a new empty line where you can type your text. The program execution is paused until you send a newline character, `\n`. Once all characters in the input have been processed the program stops again, waiting for more input. For exiting, you need to send an exit keyword that makes the `while`-loop condition false. Try to modify the program above so that:

- the program exits when you type on the space bar

**Answer:** Replace `while`-loop condition with

```
(c = getchar()) != ' ')
```

- the program exits when you send a newline character (`return`)

**Answer:** Replace `while`-loop condition with

```
(c = getchar()) != '\n')
```

- the program exits when you type `ctrl-d`

**Answer:** Replace `while`-loop condition with

```
(c = getchar()) != EOF)
```

The `ctrl-d` combination is a terminal shortcut for sending an *end of file* signal. In C, the end-of-file signal is represented by an `int`, called `EOF`, and quite often equal to `-1`, a value that is not taken by any *valid char*. Add a few lines to your code to check that `EOF = -1` on your machine. In principle, you should be careful with comparing variables of type `char` to `EOF`, as the latter is defined as an `int`. We suggest you keep this in mind and have a look at Section 1.5.1 of [The C Programming Language](#) for a discussion about `EOF` and `getchar()`. The easiest solution is to declare the variable used to store the output of `getchar()` as an `int`, i.e. to replace Line 3 with

```
int c;
```

1

**Answer:** Add these two lines to print the value of EOF

```
int i = EOF;
printf("i=%d\n", i);
```

1

2

## 2.2 Lower and upper cases

In the ASCII characters encoding, upper-case letters are ordered alphabetically from A to Z and followed by all lower-case letters, which are also ordered alphabetically from a to z, i.e.

..., A, B, ..., Z, a, b, ..., z, ...

This fact can be exploited for converting upper-case letters into lower-case letters and *vice versa*. The size of the alphabet can also be computed by subtracting the value associated with A to the value associated a, e.g. through

```
int sizeOfAlphabet;
sizeOfAlphabet = 'a' - 'A';
```

1

2

Write a function, `int upper(int c){ ... }`, that checks if the input character, `c`, is a lower case letter and, in that case, transforms it into the corresponding upper case letter. `upper` can be a modified version of

```
int lower(int c) {
    if (c >= 'A' && c <= 'Z')
        return c + 'a' - 'A';
    else
        return c;
}
```

1

2

3

4

5

6

To see the effect of `lower`, replace `putchar(c)` with `putchar(lower(c))` in `inputOutput.c`.

**Answer:**

```
#include <stdio.h>
int upper(int c) {
    if (c >= 'a' && c <= 'z')
        return c - ('a' - 'A');
    else
        return c;
}
int main() {
    int c;
    while ((c = getchar()) != -1)
        putchar(upper(c));
}
```

1

2

3

4

5

6

7

8

9

10

11

12

Finally, set the exit keyword of `inputOutput.c` to EOF and recompile it. Copy the following text

```
one two three
four five
six
```

1

2

3

into a file called `someText.txt` and observe what happens when you run

```
./a.out < someText.txt
```

1

**Answer:** The output is

```
ONE TWO THREE
FOUR FIVE
SIX
```

1

2

3

because EOF is sent automatically at the end of the text file.