

# Operating Systems Lab - Week 4: exercise

## - with answers

This lab is about low-level input-output and processes. You will start to see why studying C is important from an OS perspective. You will practice with IO process control facilities. In particular, this exercise asks you to write

- a program that reads and writes files using `stdio.h` functions for *formatted IO* and
- a program that creates a given number of child processes using `fork`.

Try to reproduce the formatted output shown in the examples *exactly*, e.g. pay attention to all capitalization details and empty spaces.

## 1 Input-output

Write a program, called `inputOutput.c`, that

- takes two file names, e.g. `fileIn.txt` and `fileOut.txt`, as *command-line* arguments,
- copies what the user writes on the terminal after the program has started into the first file, `fileIn.txt`, and
- makes a *capitalised version* of the text saved in `fileIn.txt` into the second file, `fileOut.txt`.

### 1.1 Command line inputs

To make your program accept and parse command line arguments, you need the formalism mentioned in Week 3 lab introduction. You can find an example of how an input-dependent main in the last section of Week 3's lab exercise. Try to understand what the following program does

```
#include <stdio.h>
int main(int argc, char *argv[]) {
    if (argc < 2) return -1;
    char *fileNameIn = argv[1];
    FILE *fileHandleIn = fopen(fileNameIn, "w");
    printf("fdIn=%d\n", fileno(fileHandleIn));
    printf("sizeof(fileHandleIn)=%lu\n", sizeof(*fileHandleIn));
    return 0;
}
```

How do you run the corresponding executable? Try different file names to see if the output is affected and if the file identifier, `fileno(fileHandleIn)`, changes over different runs.

**Answer:** The output is not affected because the name is stored in the struct as a string and referred to through a pointer to its first character. `fdIn` does not change over different runs. Make the program accept two files instead of one and print their identifiers and the size of their file handles. Why can't you use `printf` to print `fileHandleIn` directly?

**Answer:**

```
#include <stdio.h>
int main(int argc, char *argv[]) {
    if (argc < 3) return -1;
    char *fileNameIn = argv[1];
    char *fileNameOut = argv[2];
    FILE *fileHandleIn = fopen(fileNameIn, "w");
    FILE *fileHandleOut = fopen(fileNameOut, "w");
    printf("fdIn=%d\n", fileno(fileHandleIn));
}
```

```

printf("sizeof(fileHandleIn)=%lu\n", sizeof(*fileHandleIn));
printf("fdOut=%d\n", fileno(fileHandleOut));
printf("sizeof(fileHandleOut)=%lu\n", sizeof(*fileHandleOut));
fclose(fileHandleIn);
fclose(fileHandleOut);
return 0;
}

```

fileHandleIn is a structure of type FILE and cannot be printed with a single call of printf.

## 1.2 Parse the user input with scanf

The following program uses the `stdio.h` function `fscanf` to print a capitalized version of the user input on `stdout`.

```

#include <stdio.h>
int upper(int c) {
    if (c >= 'a' && c <= 'z')
        return c - 'a' + 'A';
    else
        return c;
}
int capitalise(char *q) {
    int c;
    int i = 0;
    while((c = *(q+i)) != '\0') {
        *(q+i) = upper(c);
        i++;
    }
    return i;
}
int main() {
    char s[10];
    while (fscanf(stdin, "%10s", s)==1) {
        capitalise(s);
        fprintf(stdout, "s=%s", s);
    }
}

```

Compile and run the program to understand how `fscanf` works. Note that

- `fscanf` is triggered by both `'\n'` and `' '`,
- to avoid overflow or other memory problems, you need to specify the maximum number of characters to be stored in the buffer through the format specifier `%10s`, and
- `capitalize` returns the length of the string and the changes in `s` are not discarded when it returns.

Look at [C online manual](#) to see how to use its return value to exit the `while`-loop.

**Answer:** `fscanf` returns -1 when it reads a non-valid character, e.g. EOF.

## 1.3 Write the original input on fileIn.txt

Use `scanf` and the file handle associated with the first file to write the user input on `fileIn.txt`. Note that `stdin` in the program above is a file handle and, to write on an open file, you can use

```
fprintf(fileHandle, "%s\n", q);
```

where `fileHandle` is a *pointer* to the structure of type FILE associated with the open file.

**Answer:**

```

#include <stdio.h>
int main(int argc, char *argv[]) {
    if (argc < 3) return -1;
    char *fileNameIn = argv[1];
    char *fileNameOut = argv[2];
    FILE *fileHandleIn = fopen(fileNameIn, "w");
    FILE *fileHandleOut = fopen(fileNameOut, "w");
    printf("fdIn=%d\n", fileno(fileHandleIn));
    printf("sizeof(fileHandleIn)=%lu\n", sizeof(*fileHandleIn));
    printf("fdOut=%d\n", fileno(fileHandleOut));
    printf("sizeof(fileHandleOut)=%lu\n", sizeof(*fileHandleOut));
    char q[10];
    while(fscanf(stdin, "%10s", q) == 1)
        fprintf(fileHandleIn, "%s\n", q);
    fclose(fileHandleIn);
    fclose(fileHandleOut);
}

```

## 1.4 Capitalize and copy the content of fileIn.txt into fileOut.txt

Complete the following program using the correct file handles. The completed program should behave as described at the beginning of this section.

```

#include <stdio.h>
int upper(int c) {
    if (c >= 'a' && c <= 'z')
        return c - 'a' + 'A';
    else
        return c;
}
int capitalise(char *q) {
    int c;
    int i = 0;
    while((c = *(q+i)) != '\0') {
        *(q+i) = upper(c);
        i++;
    }
    return i;
}
int main(int argc, char *argv[]) {
    if (argc < 3) return -1;
    char *fileNameIn = argv[1];
    char *fileNameOut = argv[2];
    FILE *fileIn = fopen(..., "w");
    char q[10];
    while(fscanf(..., "%10s", q) == 1) {
        fprintf(..., "%s", q);
    }
    fclose(...);
    ... = fopen(..., "r");
    FILE *fileOut = fopen(..., "w");
    while(fscanf(..., "%10s", q) == 1) {
        capitalise(q);
        fprintf(..., "%s", q);
    }
    fclose(...);
    fclose(...);
    return 0;
}

```

Note that the program does not print anything on screen. A run with the following user input.

```
one 1
Two 2
3
three four Five 4
5
6
and 7
s 8
i 9
x! 10
```

write

```
oneTwothreefourFiveandsix! 1
on fileIn.txt and
ONETWOTHREEFOURFIVEANDSIX! 1
on fileOut.txt.
```

### Answer:

```
#include <stdio.h> 1
int upper(int c) { 2
    if (c >= 'a' && c <= 'z') 3
        return c - 'a' + 'A'; 4
    else 5
        return c; 6
} 7
int capitalise(char *q) { 8
    int c; 9
    int i = 0; 10
    while((c = *(q+i)) != '\0') { 11
        *(q+i) = upper(c); 12
        i++; 13
    } 14
    return i; 15
} 16
17
int main(int argc, char *argv[]) { 18
    if (argc < 3) return -1; 19
    char *fileNameIn = argv[1]; 20
    char *fileNameOut = argv[2]; 21
    FILE *fileHandleIn = fopen(fileNameIn, "w"); 22
    char q[10]; 23
    while(fscanf(stdin, "%10s", q) == 1) 24
        fprintf(fileHandleIn, "%s", q); 25
    fclose(fileHandleIn); 26
    fileHandleIn = fopen(fileNameIn, "r"); 27
    FILE *fileHandleOut = fopen(fileNameOut, "w"); 28
    while(fscanf(fileHandleIn, "%10s", q) == 1) { 29
        capitalise(q); 30
        fprintf(fileHandleOut, "%s", q); 31
    } 32
    fclose(fileHandleIn); 33
    fclose(fileHandleOut); 34
    return 0; 35
} 36
```

## 2 fork

Write a program, `nChildren.c`, where a parent process creates  $N$  child processes through `fork`, waits for them to complete a task, and exits. We suggest you use the following *standard library* functions:

1. `int printf(const char *format, ...)` defined in `stdio.h` and described in Section 12.12 of [C online manual](#),
2. `pid_t getpid(void)` defined in `unistd.h` and described in Section 26.3 of [C online manual](#),
3. `pid_t fork(void)` defined in `unistd.h` and described in Section 26.4 of [C online manual](#),
4. `unsigned int sleep(int sec)` defined in `unistd.h` and described in Section 21.7 of [C online manual](#),
5. `pid_t wait(int *status)` defined in `sys/wait.h` and described in Section 26.6 of [C online manual](#), and
6. `int WEXITSTATUS(int status)` defined in `sys/wait.h` and described in Section 26.7 of [C online manual](#).

Do not forget to include the corresponding headers (`stdio.h`, `unistd.h`, `wait.h`), write

```
#include <stdio.h> 1
#include <unistd.h> 2
#include <wait.h> 3
```

at the very beginning of your code.

### 2.1 Command line argument

Again define `main` so that the program accept a single *character digit*,  $N$  as a command-line parameter, i.e. let `main` be

```
int main(int argc, char **argv) { 1
    int N = *argv[1] - '0'; 2
    ... 3
} 4
```

What is `**argv`, why is `int main(int argc, char **argv)` equivalent to `int main(int argc, char *argv[])`? Why can you use `int N = *argv[1] - '0';` to convert the input into an integer?

**Answer:** `char **argv` is a pointer to pointer and is equivalent to a *string array* as it is initialized with a list of constant strings (the user input). `*argv[1]` is a `char` containing a numerical character, and can be converted into an integer by subtracting the right offset, i.e. the ASCII code of `'0'`.

### 2.2 Write a task function

The task of all children consists of

- printing the process identifier on the terminal using `printf` and `getpid` and
- sleeping for  $n\%(N - 1)$  seconds using `sleep`.

Your function should not return any value, i.e. you should declare it as `void`, and accept two parameters, the sleeping time and the process *label*. You can use the following structure

```
void sleepingFunction(int sec, int j) { 1
    printf("%dth child (pid=%d) sleeps for %d sec\n", ..., ..., ...); 2
    sleep(...); 3
} 4
```

**Answer:**

```

void sleepingFunction(int sec, int j) {
    printf("%dth child (pid=%d) sleeps for %d sec\n", j, getpid(), sec);
    sleep(sec);
}

```

## 2.3 Generate N children with fork

You can generate a given number of child processes with a loop. Add a return statement just after the children have performed their task to avoid an uncontrolled generation of child-of-child processes. Also, make your program print a message on the terminal when one of the children terminates using `printf` and the child's label  $j = 1, \dots, N$ . For example, you can complete and add to your main the following lines

```

pid_t pid;
for (int j = 0; j < N; j++) {
    if ((pid = fork()) ...) {
        sleepingFunction(..., ...);
        printf("%dth child exits \n", ...);
        return j + 1;
    }
}

```

where  $N$  is the integer that you get from `argv`.

**Answer:**

```

int main(int argc, char **argv) {
    int N = *argv[1] - '0';
    int K = 3;
    pid_t pid;
    for (int j = 0; j < N; j++) {
        if ((pid = fork()) == 0) {
            //printf("%d", N/(2 * (j + 1)));
            sleepingFunction(N * (j % K), j + 1);
            printf("%dth child exits \n", j + 1);
            return j + 1;
        }
    }
}

```

## 2.4 Child-parent inter-process communication

Before exiting, the parent prints on the screen the order in which the children have terminated. The parents should also wait for all  $N$  children to terminate, which can be done by calling `wait`  $N$  times. Note that `pid_t wait(int *status)` returns the process identifier of the child that terminates and writes the return value of the child that terminates at the address passes as `status` parameter. To interpret the content of that address, you can call `WEXITSTATUS`, with the value stored at that address as an input. For example,

```

int status;
pidChild = wait(&status);
pidReturnValue = WEXITSTATUS(status);

```

where `pidReturnValue` is what we have called the *child label* above.

## 2.5 Print the order of arrival

Finally, the parent should print the order of the reaped children at the very end. To do this, save the return values of `wait` and `WEXITSTATUS` into two integer arrays and print their content just before the parent process terminates using

```

for (int k = 0; k < N; k++)
    printf("%dth child(pid=%d) exited %dth \n", orderVector[k], pidVector[k], k + 1);

```

**Answer:** Here is a possible version of the final program,

```

#include <stdio.h>
#include <unistd.h>
#include <wait.h>
void sleepingFunction(int sec, int j) {
    printf("%dth child (pid=%d) sleeps for %d sec\n", j, getpid(), sec);
    sleep(sec);
}
int main(int argc, char **argv) {
    int N = *argv[1] - '0';
    int K = 3;
    pid_t pid;
    for (int j = 0; j < N; j++) {
        if ((pid = fork()) == 0) {
            //printf("%d", N/(2 * (j + 1)));
            sleepingFunction(N * (j % K), j + 1);
            printf("%dth child exits \n", j + 1);
            return j + 1;
        }
    }
    int status;
    int pidVector[N];
    int orderVector[N];
    for (int k = 0; k < N; k++) {
        pidVector[k] = wait(&status);
        orderVector[k] = WEXITSTATUS(status);
    }
    for (int k = 0; k < N; k++)
        printf("%dth child(pid=%d) exited %dth \n", orderVector[k], pidVector[k], k + 1);
}

```

**Example** If  $N = 3$  a run of the program should produce an output analogous to<sup>1</sup>

```

cim-ts-node-01$ ./a.out 3
1th child (pid=3691411) sleeps for 0 seconds
2th child (pid=3691412) sleeps for 1 seconds
1th child exits
3th child (pid=3691413) sleeps for 0 seconds
3th child exits
2th child exits
1th child(pid=3691411) exited 1th
3th child(pid=3691413) exited 2th
2th child(pid=3691412) exited 3th

```

---

<sup>1</sup>Of course, you should expect different values for the process identifiers.