

CS2850 Operating System Lab

Week 8: Pipes

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

Outline

Inter-process communication

Anonymous pipes

C implementation

Processes

A *process* is the operating system's **abstraction** of a running program.

When multiple processes run **concurrently**, each process appears to have *exclusive* use of the hardware.

The process **context** is the *state information* the process needs to run.

The OS keeps the context associated with concurrent processes **separate**.

Inter-process communication

Can two processes **communicate**?

To connect the context of two processes you need an IPC *system call*.

Example: IPC occurs when a parent process *waits* for a child to terminate (by calling `wait`).

The child sends a *termination signal* to the parent that *can be decoded* through the macro `WEXITSTATUS`

example

IPC between a child and its parent.

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>
#include <stdlib.h>
int main() {
    int status, pid;
    int N = getchar() - '0';
    int *vOut = malloc(sizeof(int) * N);
    int *vPid = malloc(sizeof(int) * N);
    pid_t pid;
    for (int j = 0; j < N; j++) {
        if (! (pid=fork())) return rand() % (1 + j);
        *(vPid + j) = wait(&status);
        *(vOut + j) = WEXITSTATUS(status);
    }
    for (int k = 0; k < N; k++) printf("%d-pid=%d returns %d \n",
        k, *(vPid + k), *(vOut+k));
    free(vOut);
    free(vPid);
}
```

More general IPC

There exist different types of IPC:

- *Signals*: **simple** messages that are not used to transfer data,
- *Anonymous pipes*: **unidirectional** data channels between *related processes*,
- *Named pipes*: more flexible pipes that are treated like a **file**,
- *Shared memory*: a memory block that can be **accessed** by multiple concurrent processes,
- ...

See [Section 6 of the Linux online manual](#) or [Wikipedia page on IPC](#) for more details about UNIX and general IPC.

Linux channels

A channel is an IPC model for **message passing**.

It has a **shareable** reference that allows more processes to access it.

Processes can access the channel to **write** a message *in* it or **read** a message *from* it.

Channels normally have **two distinct references**, one for *reading* or *writing*.

They behave like **one-way** tunnels with two *gates*: an *entry* on one side and an *exit* on the other.

Writing and reading

Processes connected through a channel know the reference of (at least) one end of the channel so that

- one of them can *push* data into the channel and
- one of them can *pull* data out of the channel

i.e. reading and writing happen at different ends.

Written data *travel* through the tunnel before they are read and arrive in the **order** they are sent (FIFO).

Data written into the channel are *buffered* **by the OS** until they are read from the other end.

Implementation

In C, you can implement 3 types of channels:

- *half-duplex UNIX pipes*: communication between related processes, e.g. a parent and a child or the children of the same parent,
- *FIFOs named pipes*: communication between two independent processes,
- *sockets*: communication between different computers,
- ...

UNIX half-duplex (*anonymous*) pipes

Anonymous pipes are the **eldest** of the IPC tools.

A **half-duplex** pipe connects processes that *share a related ancestor*, i.e.

- a parent and child or
- children of the same parent.

Defining pipes requires **kernel-level operations**, i.e. a pipe is created by making a *system call*, e.g. by calling `pipe()` in a C program.

2-way pipes can be created by opening two pipes.

Example

Pipes are used in the **UNIX shell** to connect the *output* of one process to the *input* of another one, e.g. in

```
ls | sort
```

Intuitively, you can imagine data flow through the channel from the *left* to the *right* both the `ls` and `sort` processes are *children of the same process*, i.e. the shell.


Creating a pipe: step 1

The system call `pipe()` creates two **file descriptors** associated with the *reading* and *writing* ends:

- `fd0` for *reading*,
- `fd1` for *writing*.

The new pipe initially connects a process to **itself**.

Data *written to* and *read from* the pipe travels through the **kernel**.



Step 1: code

```
#include <stdio.h>
#include <unistd.h>
#define MAXCHARS 100
int main() {
    int fd[2];
    pipe(fd);
    printf("fd[0]=%d, fd[1]=%d\n", fd[0], fd[1]);
    char buf[MAXCHARS];
    *buf = '\0';
    printf("buf = %s\n", buf);
    int nbytes = write(fd[1], "hello world", 12);//write to fd[1]
    printf("%d written to fd[1]\n", nbytes);
    nbytes = read(fd[0], buf, nbytes);
    printf("%d read from fd[0]\n", nbytes);// and read from fd[0]
    printf("buf = %s\n", buf);
    return 0;
}
```

Step 1: output

The output of the program is

```
fd[0]=3, fd[1]=4  
buf =  
12 bytes written to fd[1]  
12 bytes read from fd[0]  
buf = hello world
```

File descriptors

File descriptors, e.g. `fd[0]` and `fd[1]`, are used by *low-level* I/O functions as `read` and `write`.¹

The *standard library* functions refer to files through file *handles* or **streams**, i.e. *structures* that contain more information about the file.

Given a file descriptor, you can obtain the corresponding stream and *vice versa* through.

```
FILE *streamPointer = fdopen (f0, openMode)
int f1 = fileno(streamPointer)
```

where `f0` and `f1` are (equal) integers, `openMode` a string, e.g. `"w"` or `"r"`, and `streamPointer` a pointer to a `FILE` structure.


¹See [Slides 4](#) or [Section 13.4 of the GNU online manual](#) for more details about low-level I/O

Creating a pipe: step 2

What happens if the process generates a **child** process *after creating* the pipe?

As `fork` makes an *exact copy* of the parent process, the pipe **file descriptors** are copied from the parent into the child,

Both processes have then **access** to the pipe and can use it to *communicate*.



Step 2: code

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int copyString(char *out, char *in);
int main() {
    int fd[2], nBytes;
    pipe(fd);
    char out[100], in[100], author[100];
    if (!fork()) {//child process
        sleep(1);
        nBytes=copyString(in, "world");
        copyString(author, "child");
    }
    else { //parent process
        nBytes=copyString(in, "hello");
        copyString(author, "parent");
    }
    nBytes = write(fd[1], in, nBytes);//write to fd[1]
    printf("the %s writes to fd[1]: %s (%d)\n", author, in, nBytes);
    wait(NULL);
    nBytes = read(fd[0], out, nBytes);
    printf("the %s reads from fd[0]: %s (%d)\n",author, out, nBytes);
}
```

Step 2: copyString and program output

copyString is defined by

```
int copyString(char *out, char *in) {  
    int k = 0;  
    while (*(in + k) != '\0') {  
        *(out + k) = *(in + k);  
        k++;  
    }  
    *(out + k) = '\0';  
    return k;  
}
```

The program produces the following output

```
the parent writes to fd[1]: hello (5)  
the child writes to fd[1]: world (5)  
the child reads from fd[0]: hello (5)  
the parent reads from fd[0]: world (5)
```

Problems

Why do we get such an output?

The child process inherits both **file descriptors**, `fd[0]` and `fd[1]` both processes can write to **and** read from the channel.

While the child is sleeping the parent **writes** to the channel.

The **order** of the messages is kept in the channel (channels are **FIFO** containers) so messages written by the parent and the child may get **mixed**.


Creating a pipe: step 3

You must **decide** *who writes* and *who reads* i.e. in which direction data will travel.

The processes should *close* the **unused end** of the pipe, i.e.

- close **fd[0]** if the process will be *writing* and
- close **fd[1]** if the process will be *reading*

Example: The child *writes* and the parent *reads*.



Pseudo-code

`createPipe.c` creates a message-passing system between a parent process and its child in a few steps.

- The parent creates a **half-duplex** pipe. (3)
- The parent generates a **child** process by calling `fork`. (5)
- Both processes **close** the *unused* ends of the pipe (6 and 15).
- The child's **sends** messages to the parent by calling `writeMessages`. (9)
- The parent **starts reading** the child messages and prints them on `stdout` by calling `readMessages`. (19)
- The child sends a **signal**, e.g. a message containing the *keyword* `quit`, to tell the parent that the *last message* has been sent. (10)
- The child and the parents **close** the channel. (11 and 20)

createPipe.c (main)

```
int main() {
    int fd[2];
    pipe(fd);
    char process[MAXCHARS];
    if (!fork()) {
        close(fd[0]);
        copyString(process, "the child");
        writeMessage(fd[1], "hello", process);
        writeMessage(fd[1], "world", process);
        writeMessage(fd[1], "stop", process);
        close(fd[1]);
        return 0;
    }
    else {
        close(fd[1]);
        copyString(process, "the parent");
        int end = 0;
        while(end == 0)
            end = readMessage(fd[0], process);
        close(fd[0]);
    }
}
```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22

createPipe.c (Headers and subroutines declaration)

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#define MAXCHARS 100
int readMessage(int fd, char *reader);
int writeMessage(int fd, char *buf, char *author);
int compareString(char *s1, char *s2);
int copyString(char *in, char *out);
int stringLength(char *s);
```

createPipe.c (Read and write functions)

```
int writeMessage(int fd, char *buf, char *author) {
    int nbytes = strlen(buf);
    printf("%s writes to fd[1]: %s (%d) \n", author, buf, nbytes);
    nbytes = write(fd, buf, nbytes + 1);
    return 0;
}

int readMessage(int fd, char *reader) {
    char c = ' ';
    int i = 0;
    char buf[MAXCHARS];
    while (c != '\0') {
        read(fd, &c, 1);
        buf[i++] = c;
    }
    buf[i] = '\0';
    if (strcmp(buf, "stop") == 0) return -1;
    printf("%s read from fd[0]: %s (%d) \n", reader, buf, strlen(buf));
    return 0;
}
```


Step 3: code (string functions)

```
int copyString(char *out, char *in) {
    int k = 0;
    while (*(in + k) != '\0') {
        *(out + k) = *(in + k);
        k++;
    }
    *(out + k) = '\0';
    return k;
}

int stringLength(char *s) {
    int i = 0;
    while (*(s + i) != '\0')
        i++;
    return i;
}

int compareString(char *s1, char *s2) {
    int i = 0;
    while (*(s1 + i) != '\0' && *(s2 + i) != '\0')
        if (*(s1 + i) != *(s2 + i)) return -1;
        else i++;
    if (*(s1 + i) == *(s2 + i)) return 0;
    else return -1;
}
```

Step 3: createPipe.c (output)

```
the child writes to fd[1]: hello (5)
the child writes to fd[1]: world (5)
the child writes to fd[1]: stop (4)
the parent read from fd[0]: hello (5)
the parent read from fd[0]: world (5)
```

Note: The implementation of `strlen` makes the program print the actual size of a word, excluding the null-termination character.

Would it be possible for the parent to read *hello* *before the child writes world*?