

CS2850 Operating System Lab

Week 9: Threads

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

Outline

Threads

POSIX threads

Mutexes

Threads

Computers are **busy** machines.

Threads help them perform **several tasks** at the same time.

A single **process** can solve *different related problems* by starting multiple threads.

Threads **share** the same *global* memory, e.g. data, but have their own **stack** (*automatic variables*).

Address space

A process memory contains

- **process-wide** resources, e.g. the program instructions and global data, and
- **execution state** information, e.g. the program counters and stack.

All threads can access process-wide resources but have **private** execution state information.

POSIX threads

There are **differences** between thread packages. Here we focus on **Pthreads** (for *POSIX* threads).

The **Portable Operating System Interface (POSIX)** is a set of *types and functions* that you can call from a C program (under UNIX).

Pthreads let you *divide* a program in *sub-tasks* and execute them sequentially or **in parallel**.

Pthreads in a C program

How to *create* multiple threads in C?

In a C program, you create new threads with a POSIX function called `pthread_create`.

You also need to

- define the `task` assigned to each thread (normally a subroutine of your program) and
- call a function that `merges` the output of threads when they return.

All threads are equal

There is *no hierarchy* between running threads, e.g. no parent-child distinctions.

Each thread executes **independently**: the actual execution order is *unpredictable*.

You **synchronize** them *explicitly* by **waiting** until they return.

Handling threads with POSIX *functions* requires to

- *include* the header **pthread.h** and
- *compile* your program with a **-pthread** flag.

Creating a Pthread

You create a new thread with¹

```
int pthread_create(pthread_t *thread, const
    pthread_attr_t *attr, void *subroutine, void *arg)
```

- thread points to a *buffer* that stores the thread **identifier**,
- attr points to a struct that specifies various **features** of the new thread (write NULL for the default),
- subroutine points to the **subroutine** executed by the new thread, and
- arg points to the **parameter** of start_routine.

¹See more about pthread_create on [this page](#) of the [Linux online manual on Pthreads](#).

Terminating a Pthread

You **wait** for a thread to terminate and see its **return value** with

```
int pthread_join(pthread_t thread, void **retval)
```

- thread is the thread **identifier** and
- retval is the location where the **exit status** of the target thread is copied.

A thread **terminates** when it reaches a return statement or calls

```
void pthread_exit(void *retval)
```

Synchronization

You can ensure that one *event* (in one thread) happens **before** another event (in a concurrent thread) by **synchronizing** the two threads.

For example, you can **wait** for one thread to *return* by calling `pthread_join` in the main program.


See more about thread synchronization in Chapter 3 of *PThreads Programming* by Farrel et al..

Threads or processes?

In the following examples, we implement the same program using

- no threads,
- a single thread,
- two threads, and
- two child processes.

A process with no threads



- Text contains the **program instructions**.
- Data are the **global data** needed to run the program.
- Heap is for **dynamic** memory allocation.
- Stack is for storing **automatic variables**.
- Registers, Identity, and Resources contain the **info** needed to execute.


Example: main (simple version)

```
#include <stdio.h>
void sayGoodbye(int n1, int n2);
int brian(int n), dennis(int n);
int main(){
    int nBrian = 3, nDennis = 2;
    while(nBrian + nDennis > 0) {
        if (nBrian) nBrian = brian(nBrian);
        if (nDennis) nDennis = dennis(nDennis);
        sayGoodbye(nBrian, nDennis);
    }
}
int brian(int n) {
    for (int i = 0; i < n; i++) printf("%d-hi Dennis!\n", i + 1);
    return n - 1;
}
int dennis(int n) {
    for (int i = 0; i < n; i++) printf("%d-hi Brian!\n", i + 1);
    return n - 1;
}
void sayGoodbye(int n1, int n2) {
    printf("%d-goodbye!\n", n1 + n2);
}
```

Output

```
./a.out  
1-hi Dennis!  
2-hi Dennis!  
3-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
3-goodbye!  
1-hi Dennis!  
2-hi Dennis!  
1-hi Brian!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```

A process with a single Pthreads



Stack and Registers are now part of the thread.

Each *active thread* has a **reserved stack frame** and **specific machine registers**.

Example: main (Pthread compatible)

pthread_create takes three free arguments

```
int pthread_create(pthread_t *tName, NULL, void *f(void
    *), void *arg)
```

void *f(void *) is (a pointer to) a function of type void which takes a pointer to void as a single argument.

The function arguments should be passed through a single pointer to void.

You need to rewrite brian and dennis as

```
void *brian(void *n)
void *dennis(void *n)
```


New version of brian and dennis

To meet the requirement of `pthread_create` you need a structure holding the original integer argument,

```
struct arg {  
    int n;  
};
```

In the first line of `f`, cast the void pointer to be a pointer to a variable of type `struct arg`²

```
struct arg *p = n;
```

In main, `n` becomes `(&argName)->n`, where `argName` is an object of type `struct arg`.

²To access the original arguments with `p->n`.

New version of brian and dennis

You obtain

```
void *brian(void *nIn) {  
    struct arg *nStack = nIn;  
    for (int i=0; i<nStack->n; i++)  
        printf("%d-hi Dennis!\n", i + 1);  
    nStack->n = nStack->n - 1;  
    return NULL;  
}  
  
void *dennis(void *nIn) {  
    struct arg *nStack = nIn;  
    for (int i=0; i<nStack->n; i++)  
        printf("%d-hi Brian!\n", i + 1);  
    nStack->n = nStack->n - 1;  
    return NULL;  
}
```

New version of brian and dennis

The new version of main is

```
#include <stdio.h>

struct arg{
    int n;
};

void *brian(void *n), *dennis(void *n), sayGoodbye(int n1, int n2);

int main(){
    struct arg nBrian = {3}, nDennis = {2};
    while(nBrian.n + nDennis.n > 0) {
        if (nBrian.n) brian(&nBrian);
        if (nDennis.n) dennis(&nDennis);
        sayGoodbye(nBrian.n, nDennis.n);
    }
}
```

sayGoodbye can be left unchanged.

Example (1 running Pthread)

```
#include <stdio.h>
#include <pthread.h>

struct arg{
    int n;
};

void *brian(void *n), *dennis(void *n), sayGoodbye(int n1, int n2);

int main(){
    pthread_t t1;
    struct arg nBrian = {3}, nDennis = {2};
    while(nBrian.n + nDennis.n > 0) {
        if (nBrian.n) pthread_create(&t1, NULL, brian, &nBrian);
        pthread_join(t1, NULL);
        if (nDennis.n) pthread_create(&t1, NULL, dennis, &nDennis);
        pthread_join(t1, NULL);
        sayGoodbye(nBrian.n, nDennis.n);
    }
}
```


All auxiliary functions are defined as above.

Output (1 running Pthread)

The execution is as before

```
./a.out  
1-hi Dennis!  
2-hi Dennis!  
3-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
3-goodbye!  
1-hi Dennis!  
2-hi Dennis!  
1-hi Brian!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```

A process with two threads



Both threads have *their private copy* of the machine **registers** and can use variables or file descriptors in the *process-wide areas*.

The two threads execute at **different locations**.

Example (2 running Pthreads)

```
#include <stdio.h>
#include <pthread.h>


struct arg{
    int n;
};

void *brian(void *n), *dennis(void *n), sayGoodbye(int n1, int n2);

int main(){
    struct arg nBrian = {3}, nDennis = {2};
    pthread_t t1, t2;
    while(nBrian.n + nDennis.n > 0) {
        if (nBrian.n) pthread_create(&t1, NULL, brian, &nBrian);
        if (nDennis.n) pthread_create(&t2, NULL, dennis, &nDennis);
        pthread_join(t1, NULL);
        pthread_join(t2, NULL);
        sayGoodbye(nBrian.n, nDennis.n);
    }
}
```

As t1 starts before t2 and the task performed by brian is *computationally cheap*, the output on stdout is unchanged.

Two processes vs. two threads



The address space of two concurrent processes created by `fork` is *completely private* (no shared areas).

The two processes can interact via a **channel** (anonymous pipe).

Example: main (2 running child processes, slowBrian)

```
int main() {  
    int status;  
    while(nBrian.n + nDennis.n > 0) {  
        int pidBrian = fork();  
        if (!pidBrian) {  
            if (nBrian.n) brianSlow(&nBrian);  
            return nBrian.n;  
        }  
        int pidDennis = fork();  
        if (!pidDennis) {  
            if (nDennis.n) dennis(&nDennis);  
            return nDennis.n;  
        }  
        waitpid(pidBrian, &status, 0);  
        nBrian.n = WEXITSTATUS(status);  
        waitpid(pidDennis, &status, 0);  
        nDennis.n = WEXITSTATUS(status);  
        sayGoodbye(nBrian.n, nDennis.n);  
    }  
}
```

Include `unistd.h` and `sys/wait.h` to use `fork` and `wait`.

Output (2 running child processes, slow Brian)

The output is the same as before

```
./a.out  
1-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
2-hi Dennis!  
3-hi Dennis!  
3-goodbye!  
1-hi Dennis!  
1-hi Brian!  
2-hi Dennis!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```

Another example of synchronization: Mutexes

Variables of type `pthread_mutex_t` allow you to synchronize the **access** to some *area of your code*.

The protected area is **mutually exclusive**. While one thread is in the protected area, other threads cannot

- *execute* the protected code and
- *lock or unlock* the Mutex.

Mutex-protected variables

To use a Mutex to regulate the *updates of shared variable*, *x*,

1. create and initialize a **mutex variable**, *m*, using

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;  
(m should be declared outside main).
```

2. in the thread subroutine, *lock and unlock* the mutex before updating *x*, e.g. write

```
pthread_mutex_lock(&m)  
x= ...  
pthread_mutex_unlock(&m)
```

Example: (2 running threads with thread-safe subroutines)

The *thread-safe* version of slowBrian and dennis is

```
void *brianSlow(void *n) {
    struct arg *nl = n;
    int t = 0;
    pthread_mutex_lock(&m); //start of the mutual exclusive area
    for (int i=0; i<(nl->n); i++) {
        printf("%d-hi Dennis!\n", i + 1);
        while (t < 100000) t++;
    }
    pthread_mutex_unlock(&m); //end of the mutual exclusive area
    nl->n = nl->n - 1;
    return NULL;
}

void *dennis(void *n) {
    struct arg *nl = n;
    pthread_mutex_lock(&m); //start of the mutual exclusive area
    for (int i=0; i<(nl->n); i++) printf("%d-hi Brian!\n", i + 1);
    pthread_mutex_unlock(&m); //end of the mutual exclusive area
    (nl->n) = (nl->n) - 1;
    return NULL;
}
```

Output (2 running threads, thread-safe subroutines)

```
./a.out  
1-hi Dennis!  
2-hi Dennis!  
3-hi Dennis!  
1-hi Brian!  
2-hi Brian!  
3-goodbye!  
1-hi Dennis!  
2-hi Dennis!  
1-hi Brian!  
1-goodbye!  
1-hi Dennis!  
0-goodbye!
```