

CS2850 Operating System Lab

Week 10: Review of pointers, strings, and process control

nicolo colombo

`nicolo.colombo@rhul.ac.uk`

Office Bedford 2-21

outline

Pointers

Strings

Pointers and functions

fork

Pointers (1)

Pointers are **variables** that store the memory address of other variables.

Address operator: `&i` is the **address** of `i`.

Dereferencing operator: `*ip` is the **value** of the variable stored at the address `ip`.

You can **print** on stdout the *unpredictable value* of `ip`, i.e. the address stored in `ip`, with

```
printf("%p", ( void *) ip)
```

Pointers (2)

The **type** of a pointer is the type of variable stored at the pointed address.

The data type of a variable specifies how to **interpret** the data stored in the variable.

A pointer of a given type **cannot** store the address of variables of a different type, e.g.

```
int i = 1;  
char *pc = &i
```

is **illegal** and produces compilation errors.

Example: int, char, int*, and char* (1)

```
#include <stdio.h>
int main() {
    int i = 10;
    int *pi = &i;
    char c = 'a';
    char *pc = &c;
    printf("sizeof(pi)=%lu\n", sizeof(pi));
    printf("sizeof(pc)=%lu\n", sizeof(pc));
    printf("pi=%p and  &i = %p\n", (void *) pi, (void *) &
        i);
    printf("pc=%p and &c = %p\n", (void *) pc, (void *) &c
        );
    printf("*pi=%d and i = %d\n", *pi, i);
    printf("*pc=%c and c = %d\n", *pc, c);
}
```

Output

```
./a.out
sizeof(pi)=8
sizeof(pc)=8
pi=0x7ffd17e6b974 and &i = 0x7ffd17e6b974
pc=0x7ffd17e6b973 and &c = 0x7ffd17e6b973
*pi=10 and i = 10
*pc=a and c = 97
```

Strings (1)

String constants are *null-terminated* arrays of char, e.g.

```
char *s = "hello world"
```

s is a pointer to char that stores the address of the first element of the string "hello world".

The last char of s is the null character '**\0**'.

Strings (2)

String *constants* and character arrays, e.g. `s` and `as` defined by

```
char *s = "hello world";  
char as[12] = "hello world";
```

are *not equivalent*.

`s` is a pointer (a variable) and can be *assigned to a different address*.

`as` is an array and always refers to *fixed allocated* storage.

Strings (3)

`s[7]` and `as[7]` both return `'w'` but

```
s[7] = 'x';
```

causes a Segmentation Fault, while

```
as[7] = 'x';
```

is **legal** and sets the 8th character of `as` to `'x'`, i.e. `'o'` becomes `'x'`.

Strings (4)

You can use `printf("%s", s)` and `printf("%s", as)` to **print** the strings `s` and `as` **up to** their null-termination character.

In C, there are *no built-in operators* for processing an **entire** string as a unit, e.g.

```
char as[12];  
as = 'hi world';
```

is *illegal*, but you can *re-assign* `s` to point to another string constant, e.g.

```
char *s;  
s = 'hi world';
```

Strings (6)

You can print **a portion** of `s` or `as` by specifying the address of a single character, e.g.

```
s = "hi world";  
printf("(s + 3) = %s", (s + 3));
```

prints (s + 3) = world.

You can **cut** a character array by setting one element to `'\0'`, e.g.

```
char as[10] = "hi world";  
*(as + 2) = '\0';  
printf("as = %s and (as + 3) = %s\n", s, (s + 3));
```

prints as = hi and (as + 3) = world.

Example: char* and char[13]

```
#include <stdio.h>
int main() {
    char *s = "hello world";
    char sa[13] = "hello world";
    //s[7] = 'x'; is illegal
    sa[7] = 'x'; //is legal
    printf("s[7] = %c\n", s[7]);
    printf("sa[7] = %c\n", sa[7]);
    printf("(s + 1) = %s\n", s + 1) ;
    printf("(sa + 1) = %s\n", sa + 1);
    s = "hi world";
    printf("s = %s\n", s);
    //sa = s; is illegal
    s = &sa[6]; //is legal and equivalent to s = sa + 6;
    printf("s = &sa[6] => s=%s\n", s);
    printf("(sa + 6) = %p and s = %p\n", (void *) (sa + 6),
        (void *) s);
}
```

Output

```
./a.out  
s[7] = o  
sa[7] = x  
(s + 1) = ello world  
(sa + 1) = ello wxrld  
s = hi world  
s = &sa[6] => s=wxrld  
(sa + 6) = 0x7ffd23e608c1 and s = 0x7ffd23e608c1
```

Example: pointer arithmetic

```
#include <stdio.h>
int main() {
    char as[100] = "one two three four five";
    char vs[100] = "one two three four five";
    char *s = "one two three four five";
    for (int i = 0; i < 13; i++) s++;
    as[13] = '\\0';
    printf("as = %s\\n", as);
    printf("vs = %s\\n", vs);
    printf("s = %s\\n", s);
}
```

Output:

```
as = one two three
vs = one two three four five
s = four five
```

Pointers and functions (1)

The de-referencing operator can be used in **function declarations**, e.g.

```
double f(char *c);
```

says that f **returns** a double and the **argument** of f is a pointer to char.

Arguments are passed to functions **by value**. There is *no direct way* for the called function to **modify** a variable in the calling function.

To change the value of a variable in the calling function, pass **the address** of that variable.

Pointers and functions (2)

Arrays are represented by *the location of the initial element*.

When you pass an array to a function, you will see the changes in `main`, e.g.

```
#include <stdio.h>
void initialise(int *a) {
    int i;
    for (i = 0; i < 10; i++) *(a + i) = i;
}
int main() {
    int i, a[10];
    initialise(a);
    for (i = 0; i < 10; i++) printf("%d ", *(a + i));
}
```

prints 0 1 2 3 4 5 6 7 8 9.

Example: swapping int

```
#include <stdio.h>
void swap(int *pi, int *pj);
void swapWrong(int i, int j);
int main() {
    int x = 1, y = 2, *px = &x, *py = &y;
    printf("x = %d, y = %d, px = %p, and py = %p\n", x, y, (void *) px, (void *) py);
    printf("calling swapWrong ... \n");
    swapWrong(x, y);
    printf("x = %d, y = %d, px = %p, and py = %p\n", x, y, (void *) px, (void *) py);
    printf("calling swap ... \n");
    swap(&x, &y);
    printf("x = %d, y = %d, px = %p, and py = %p\n", x, y, (void *) px, (void *) py);
}
```

Output:

```
./a.out
x = 1, y = 2, px = 0x7fff73496720, and py = 0x7fff73496724
calling swapWrong ...
x = 1, y = 2, px = 0x7fff73496720, and py = 0x7fff73496724
calling swap ...
x = 2, y = 1, px = 0x7fff73496720, and py = 0x7fff73496724
```

Auxiliary functions

A function to swap the *content of two addresses*.

```
void swap(int *pi, int *pj) {  
    int temp = *pi;  
    *pi = *pj;  
    *pj = temp;  
}
```

A function to swap two *stack variables*.¹

```
void swapWrong(int i, int j) {  
    int temp = i;  
    i = j;  
    j = temp;  
}
```

¹All changes are discarded when the function returns.

Process control (1)

`unistd.h` and `sys/wait.h` contains the functions you need to **control** a process from a C program.

`int getpid()` returns the **process identifier** (PID) of the running process.

`int fork()` creates a **child** process that is a copy of the running process.

`void exit()` or return **terminate** the current process.

`int wait(int * status):` to **wait** for the *first* child that terminates (call it several times if you have generated more than one child). ²

²See more about process control in [Section 26 of the GNU online manual](#).

Process control (2)

The OS assigns PIDs in **increasing order** but their values are unpredictable.

In C, `fork` is *called once but returns twice*.

The return value is -1 if the process creation failed, 0 in the child process, and the child's PID in the parent process.

Process control (3)

When `fork` returns, the child address space is a copy of the parent address space, but the parent and the child have **different PIDs**.

The **return values** of `fork` allow you to make conditional statements.

The address spaces of the child and the parent are *private*, i.e. the parent cannot see the **changes** made by the child, and vice versa.

The parent and the child **run concurrently**. The execution order is unpredictable.

Example: private but identical address spaces

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main() {
    int j = 0;
    printf("PID = %d, j = %d, &j = %p \n", getpid(), j, (void *) &j);
    int pid = fork();
    if (!pid) j = 1;
    else j = 2;
    printf("PID = %d, j = %d, &j = %p \n", getpid(), j, (void *) &j);
}
```

Output:

```
./a.out
PID = 2789388, j = 0, &j = 0x7fff5ea69010
PID = 2789388, j = 2, &j = 0x7fff5ea69010
PID = 2789389, j = 1, &j = 0x7fff5ea69010
```